

# Predicting House Sale Prices in King County, WA.

2023-12-14

```
library(openxlsx)
```

```
fp <- "C:\\Users\\PC\\Downloads\\kc_house_data.xlsx"
```

```
hp_data <- read.xlsx(fp)
```

```
head(hp_data)
```

```
##           id           date    price bedrooms bathrooms sqft_living sqft_lot
## 1 7129300520 20141013T000000  221900         3         1.00        1180     5650
## 2 6414100192 20141209T000000  538000         3         2.25        2570     7242
## 3 5631500400 20150225T000000  180000         2         1.00         770    10000
## 4 2487200875 20141209T000000  604000         4         3.00        1960     5000
## 5 1954400510 20150218T000000  510000         3         2.00        1680     8080
## 6 7237550310 20140512T000000 1230000         4         4.50        5420    101930
##  floors waterfront view condition grade sqft_above sqft_basement yr_built
## 1         1           0     0           3       7        1180           0     1955
## 2         2           0     0           3       7        2170          400     1951
## 3         1           0     0           3       6         770           0     1933
## 4         1           0     0           5       7        1050          910     1965
## 5         1           0     0           3       8        1680           0     1987
## 6         1           0     0           3      11        3890        1530     2001
##  yr_renovated zipcode      lat      long sqft_living15 sqft_lot15
## 1           0     98178 47.5112 -122.257        1340        5650
## 2          1991     98125 47.7210 -122.319        1690        7639
## 3           0     98028 47.7379 -122.233        2720        8062
## 4           0     98136 47.5208 -122.393        1360        5000
## 5           0     98074 47.6168 -122.045        1800        7503
## 6           0     98053 47.6561 -122.005        4760       101930
```

```
str(hp_data)
```

```
## 'data.frame':    21613 obs. of  21 variables:
##  $ id           : num  7.13e+09 6.41e+09 5.63e+09 2.49e+09 1.95e+09 ...
##  $ date          : chr   "20141013T000000" "20141209T000000" "20150225T000000" "20141209T000000" ...
##  $ price         : num  221900 538000 180000 604000 510000 ...
##  $ bedrooms      : num  3 3 2 4 3 4 3 3 3 3 ...
##  $ bathrooms     : num  1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
##  $ sqft_living   : num  1180 2570 770 1960 1680 ...
##  $ sqft_lot      : num  5650 7242 10000 5000 8080 ...
##  $ floors        : num  1 2 1 1 1 1 2 1 1 2 ...
##  $ waterfront    : num  0 0 0 0 0 0 0 0 0 0 ...
##  $ view          : num  0 0 0 0 0 0 0 0 0 0 ...
##  $ condition     : num  3 3 3 5 3 3 3 3 3 3 ...
```

```
## $ grade      : num 7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above : num 1180 2170 770 1050 1680 ...
## $ sqft_basement: num 0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built    : num 1955 1951 1933 1965 1987 ...
## $ yr_renovated : num 0 1991 0 0 0 ...
## $ zipcode     : num 98178 98125 98028 98136 98074 ...
## $ lat         : num 47.5 47.7 47.7 47.5 47.6 ...
## $ long        : num -122 -122 -122 -122 -122 ...
## $ sqft_living15: num 1340 1690 2720 1360 1800 ...
## $ sqft_lot15  : num 5650 7639 8062 5000 7503 ...
```

*#Checking for missing values*

```
mv <- apply(hp_data, function(x) sum(is.na(x)))
print(mv)
```

```
##      id      date      price      bedrooms      bathrooms
##      0      0      0      0      0
## sqft_living sqft_lot      floors      waterfront      view
##      0      0      0      0      0
##      condition      grade      sqft_above sqft_basement      yr_built
##      0      0      0      0      0
## yr_renovated      zipcode      lat      long sqft_living15
##      0      0      0      0      0
##      sqft_lot15
##      0
```

*#checking for duplicates*

```
dup_values <- sum(duplicated(hp_data))
print(dup_values)
```

```
## [1] 0
```

*#checking for negative values*

```
negatives_hp_data <- apply(hp_data, function(col) any(col < 0))
print(negatives_hp_data)
```

```
##      id      date      price      bedrooms      bathrooms
##      FALSE      FALSE      FALSE      FALSE      FALSE
## sqft_living sqft_lot      floors      waterfront      view
##      FALSE      FALSE      FALSE      FALSE      FALSE
##      condition      grade      sqft_above sqft_basement      yr_built
##      FALSE      FALSE      FALSE      FALSE      FALSE
## yr_renovated      zipcode      lat      long sqft_living15
##      FALSE      FALSE      FALSE      TRUE      FALSE
##      sqft_lot15
##      FALSE
```

```
negative_rows_long <- sum(hp_data$long < 0)
print(negative_rows_long)
```

```
## [1] 21613
```

```
str(hp_data)
```

```
## 'data.frame': 21613 obs. of 21 variables:
## $ id : num 7.13e+09 6.41e+09 5.63e+09 2.49e+09 1.95e+09 ...
## $ date : chr "20141013T000000" "20141209T000000" "20150225T000000" "20141209T000000" ...
## $ price : num 221900 538000 180000 604000 510000 ...
## $ bedrooms : num 3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms : num 1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_living : num 1180 2570 770 1960 1680 ...
## $ sqft_lot : num 5650 7242 10000 5000 8080 ...
## $ floors : num 1 2 1 1 1 1 2 1 1 2 ...
## $ waterfront : num 0 0 0 0 0 0 0 0 0 0 ...
## $ view : num 0 0 0 0 0 0 0 0 0 0 ...
## $ condition : num 3 3 3 5 3 3 3 3 3 3 ...
## $ grade : num 7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above : num 1180 2170 770 1050 1680 ...
## $ sqft_basement: num 0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built : num 1955 1951 1933 1965 1987 ...
## $ yr_renovated : num 0 1991 0 0 0 ...
## $ zipcode : num 98178 98125 98028 98136 98074 ...
## $ lat : num 47.5 47.7 47.7 47.5 47.6 ...
## $ long : num -122 -122 -122 -122 -122 ...
## $ sqft_living15: num 1340 1690 2720 1360 1800 ...
## $ sqft_lot15 : num 5650 7639 8062 5000 7503 ...
```

```
#removing the id column
hp_data_1 <- hp_data[, -(1)]
str(hp_data_1)
```

```
## 'data.frame': 21613 obs. of 20 variables:
## $ date : chr "20141013T000000" "20141209T000000" "20150225T000000" "20141209T000000" ...
## $ price : num 221900 538000 180000 604000 510000 ...
## $ bedrooms : num 3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms : num 1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_living : num 1180 2570 770 1960 1680 ...
## $ sqft_lot : num 5650 7242 10000 5000 8080 ...
## $ floors : num 1 2 1 1 1 1 2 1 1 2 ...
## $ waterfront : num 0 0 0 0 0 0 0 0 0 0 ...
## $ view : num 0 0 0 0 0 0 0 0 0 0 ...
## $ condition : num 3 3 3 5 3 3 3 3 3 3 ...
## $ grade : num 7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above : num 1180 2170 770 1050 1680 ...
## $ sqft_basement: num 0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built : num 1955 1951 1933 1965 1987 ...
## $ yr_renovated : num 0 1991 0 0 0 ...
## $ zipcode : num 98178 98125 98028 98136 98074 ...
## $ lat : num 47.5 47.7 47.7 47.5 47.6 ...
## $ long : num -122 -122 -122 -122 -122 ...
## $ sqft_living15: num 1340 1690 2720 1360 1800 ...
## $ sqft_lot15 : num 5650 7639 8062 5000 7503 ...
```

```
head(hp_data_1)
```

```
##           date    price bedrooms bathrooms sqft_living sqft_lot floors
## 1 20141013T000000 221900         3         1.00        1180    5650      1
## 2 20141209T000000 538000         3         2.25        2570    7242      2
## 3 20150225T000000 180000         2         1.00         770   10000      1
## 4 20141209T000000 604000         4         3.00        1960    5000      1
## 5 20150218T000000 510000         3         2.00        1680    8080      1
## 6 20140512T000000 1230000        4         4.50        5420   101930     1
##   waterfront view condition grade sqft_above sqft_basement yr_built
## 1          0    0          3     7        1180           0    1955
## 2          0    0          3     7        2170          400    1951
## 3          0    0          3     6         770           0    1933
## 4          0    0          5     7        1050          910    1965
## 5          0    0          3     8        1680           0    1987
## 6          0    0          3    11        3890         1530    2001
##   yr_renovated zipcode      lat      long sqft_living15 sqft_lot15
## 1            0   98178 47.5112 -122.257        1340        5650
## 2          1991   98125 47.7210 -122.319        1690        7639
## 3            0   98028 47.7379 -122.233        2720        8062
## 4            0   98136 47.5208 -122.393        1360        5000
## 5            0   98074 47.6168 -122.045        1800        7503
## 6            0   98053 47.6561 -122.005        4760       101930
```

```
# handling the date column
```

```
hp_data_1$date <- as.Date(hp_data_1$date , format="%Y%m%d")
```

```
#splitting the date column
```

```
hp_data_1$year <- as.numeric(format(hp_data_1$date, "%Y"))
```

```
hp_data_1$month <- as.numeric(format(hp_data_1$date, "%m"))
```

```
head(hp_data_1)
```

```
##           date    price bedrooms bathrooms sqft_living sqft_lot floors waterfront
## 1 2014-10-13 221900         3         1.00        1180    5650      1          0
## 2 2014-12-09 538000         3         2.25        2570    7242      2          0
## 3 2015-02-25 180000         2         1.00         770   10000      1          0
## 4 2014-12-09 604000         4         3.00        1960    5000      1          0
## 5 2015-02-18 510000         3         2.00        1680    8080      1          0
## 6 2014-05-12 1230000        4         4.50        5420   101930     1          0
##   view condition grade sqft_above sqft_basement yr_built yr_renovated zipcode
## 1    0          3     7        1180           0    1955           0   98178
## 2    0          3     7        2170          400    1951         1991   98125
## 3    0          3     6         770           0    1933           0   98028
## 4    0          5     7        1050          910    1965           0   98136
## 5    0          3     8        1680           0    1987           0   98074
## 6    0          3    11        3890         1530    2001           0   98053
##           lat      long sqft_living15 sqft_lot15 year month
## 1 47.5112 -122.257        1340        5650 2014    10
## 2 47.7210 -122.319        1690        7639 2014    12
## 3 47.7379 -122.233        2720        8062 2015     2
## 4 47.5208 -122.393        1360        5000 2014    12
## 5 47.6168 -122.045        1800        7503 2015     2
## 6 47.6561 -122.005        4760       101930 2014     5
```

```
#removing the date column
hp_data_2 <- hp_data_1[, -(1)]
head(hp_data_2)
```

```
##      price bedrooms bathrooms sqft_living sqft_lot floors waterfront view
## 1  221900         3       1.00        1180    5650      1          0    0
## 2  538000         3       2.25         2570    7242      2          0    0
## 3  180000         2       1.00          770   10000      1          0    0
## 4  604000         4       3.00         1960    5000      1          0    0
## 5  510000         3       2.00         1680    8080      1          0    0
## 6 1230000         4       4.50         5420   101930      1          0    0
##      condition grade sqft_above sqft_basement yr_built yr_renovated zipcode
## 1           3      7        1180           0     1955          0  98178
## 2           3      7        2170          400     1951         1991  98125
## 3           3      6         770           0     1933          0  98028
## 4           5      7        1050          910     1965          0  98136
## 5           3      8         1680           0     1987          0  98074
## 6           3     11         3890         1530     2001          0  98053
##      lat      long sqft_living15 sqft_lot15 year month
## 1 47.5112 -122.257        1340        5650 2014    10
## 2 47.7210 -122.319        1690        7639 2014    12
## 3 47.7379 -122.233        2720        8062 2015     2
## 4 47.5208 -122.393        1360        5000 2014    12
## 5 47.6168 -122.045        1800        7503 2015     2
## 6 47.6561 -122.005        4760       101930 2014     5
```

```
str(hp_data_2)
```

```
## 'data.frame':  21613 obs. of  21 variables:
## $ price      : num  221900 538000 180000 604000 510000 ...
## $ bedrooms   : num  3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms   : num  1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_living : num  1180 2570 770 1960 1680 ...
## $ sqft_lot    : num  5650 7242 10000 5000 8080 ...
## $ floors      : num  1 2 1 1 1 1 2 1 1 2 ...
## $ waterfront  : num  0 0 0 0 0 0 0 0 0 0 ...
## $ view        : num  0 0 0 0 0 0 0 0 0 0 ...
## $ condition   : num  3 3 3 5 3 3 3 3 3 3 ...
## $ grade       : num  7 7 6 7 8 11 7 7 7 7 ...
## $ sqft_above  : num  1180 2170 770 1050 1680 ...
## $ sqft_basement: num  0 400 0 910 0 1530 0 0 730 0 ...
## $ yr_built    : num  1955 1951 1933 1965 1987 ...
## $ yr_renovated: num  0 1991 0 0 0 ...
## $ zipcode     : num  98178 98125 98028 98136 98074 ...
## $ lat         : num  47.5 47.7 47.7 47.5 47.6 ...
## $ long        : num  -122 -122 -122 -122 -122 ...
## $ sqft_living15: num  1340 1690 2720 1360 1800 ...
## $ sqft_lot15  : num  5650 7639 8062 5000 7503 ...
## $ year        : num  2014 2014 2015 2014 2015 ...
## $ month       : num  10 12 2 12 2 5 6 1 4 3 ...
```

```

#Picking out the continuous predictors
continuous_vars <- c("bedrooms", "bathrooms", "sqft_living", "sqft_lot", "floors", "sqft_above", "sqft_living15", "yr_built", "yr_renovated", "lat", "long", "sqft_living15", "sqft_lot15")

df_continuous <- hp_data_2[, continuous_vars]

#Boxplot

# Names of continuous predictors
predictor_names <- names(df_continuous)

# Function to generate boxplots for all continuous predictors
create_boxplots <- function(data, predictor_names) {
  # Create a layout for the boxplots
  par(mfrow = c(3, 3)) # Change the rows and columns as needed

  # Loop through each predictor and create a boxplot
  for (predictor in predictor_names) {
    boxplot(data[[predictor]],
            main = paste("Boxplot of", predictor),
            xlab = predictor,
            col = "lightblue",
            border = "blue",
            notch = TRUE)
  }
}

# Function to create boxplots for all predictors
create_boxplots(df_continuous, predictor_names)

```

```

## Warning in (function (z, notch = FALSE, width = NULL, varwidth = FALSE, : some
## notches went outside hinges ('box'): maybe set notch=FALSE

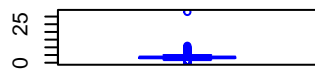
```

```

## Warning in (function (z, notch = FALSE, width = NULL, varwidth = FALSE, : some
## notches went outside hinges ('box'): maybe set notch=FALSE

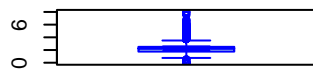
```

**Boxplot of bedrooms**



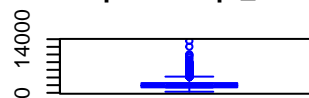
bedrooms

**Boxplot of bathrooms**



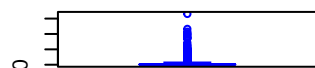
bathrooms

**Boxplot of sqft\_living**



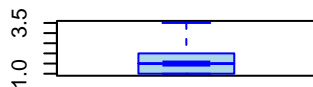
sqft\_living

**Boxplot of sqft\_lot**



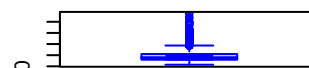
sqft\_lot

**Boxplot of floors**



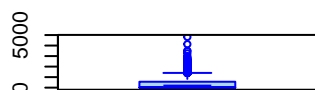
floors

**Boxplot of sqft\_above**



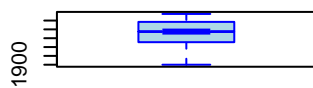
sqft\_above

**Boxplot of sqft\_basement**



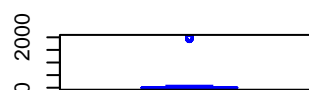
sqft\_basement

**Boxplot of yr\_built**

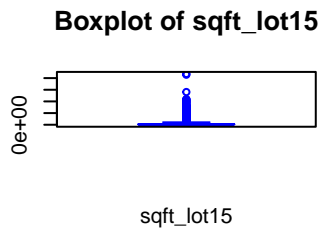
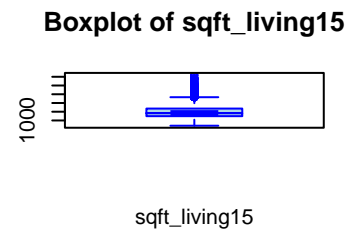
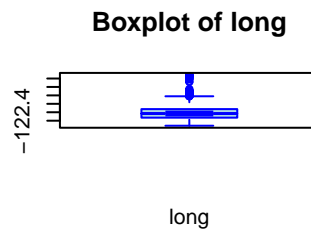
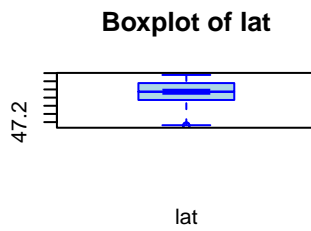


yr\_built

**Boxplot of yr\_renovated**



yr\_renovated



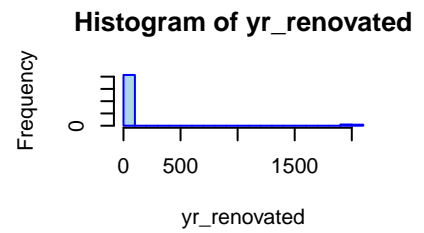
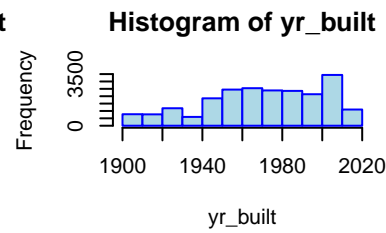
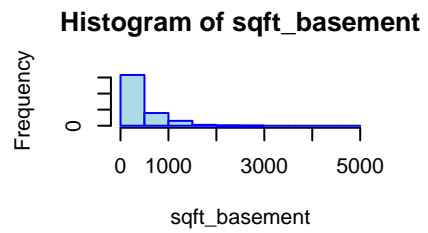
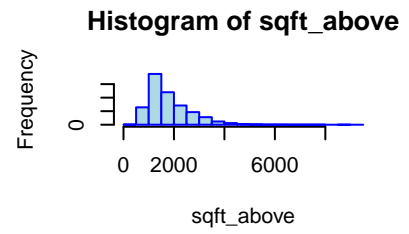
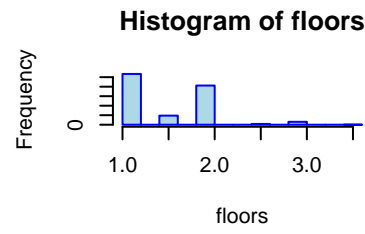
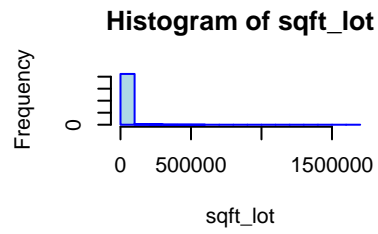
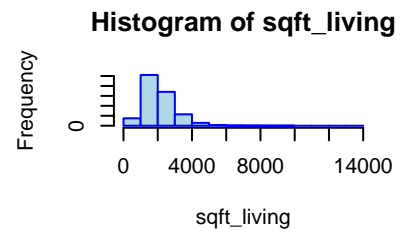
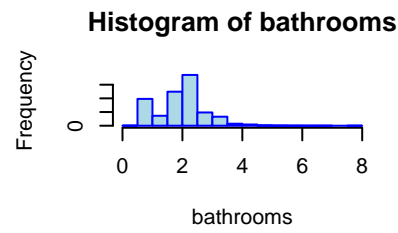
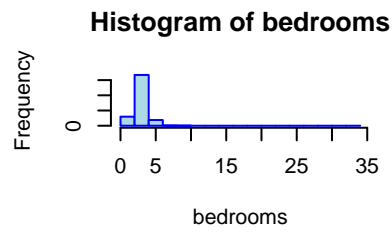
```
#Histogram

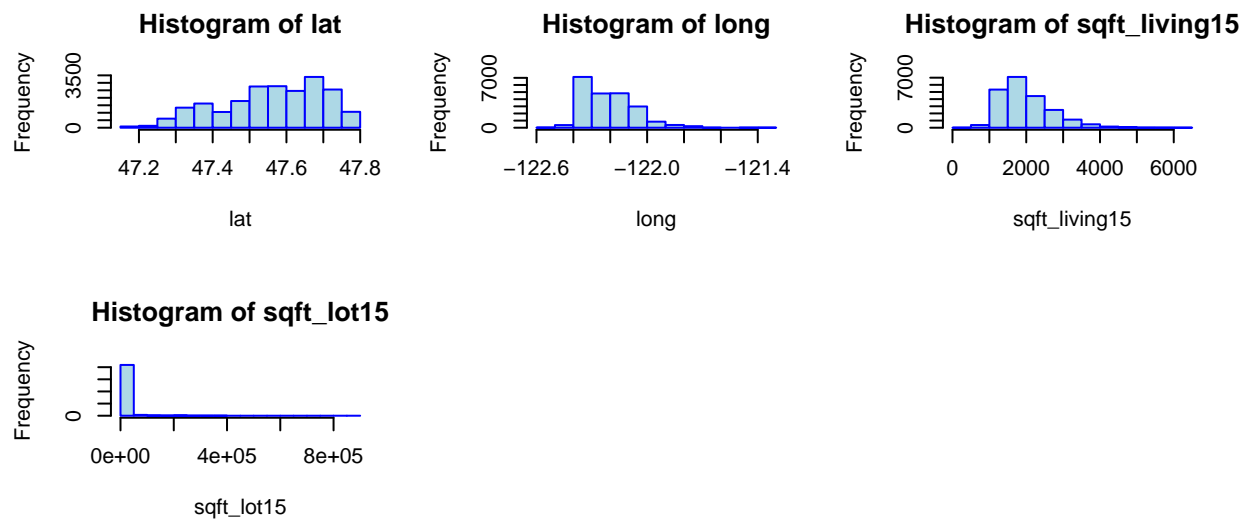
# Function to generate histograms for continuous predictors
create_histograms <- function(data, predictor_names) {
  # Create a layout for the histograms
  par(mfrow = c(3, 3)) # Change the rows and columns as needed

  # Loop through each predictor and create a histogram
  for (predictor in predictor_names) {
    hist(data[[predictor]],
          main = paste("Histogram of", predictor),
          xlab = predictor,
          col = "lightblue",
          border = "blue")
  }
}

# Function to create histograms for continuous predictors
create_histograms(df_continuous, predictor_names)
```



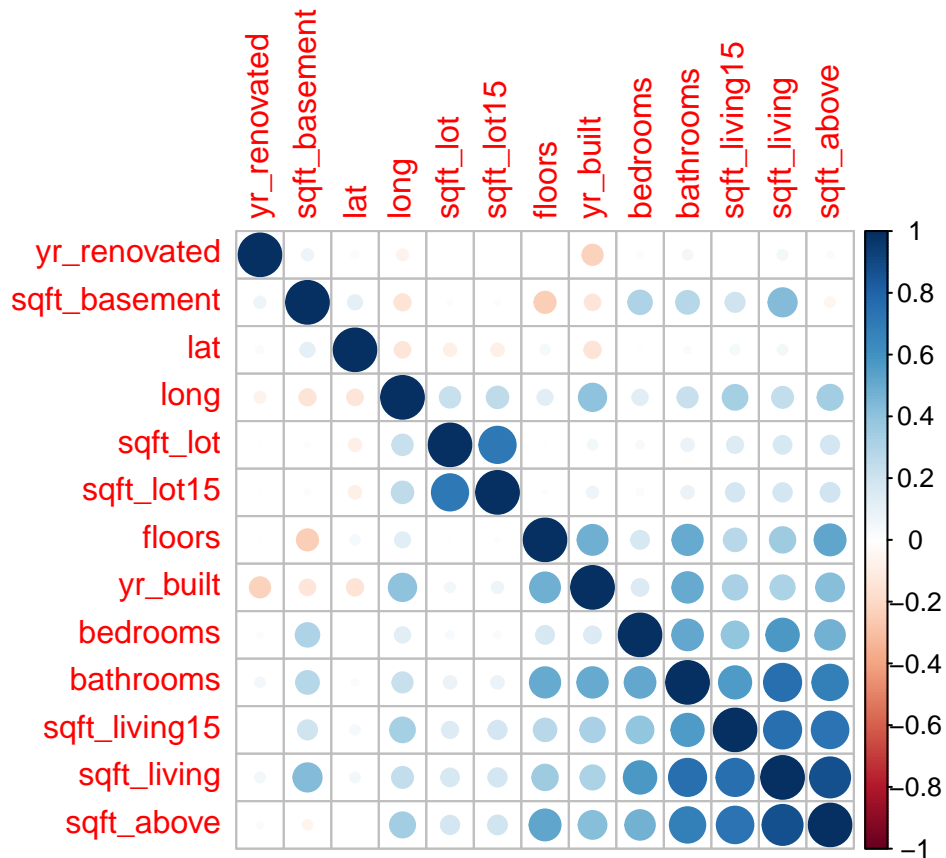




```
#####corr plot
correlations <- cor(df_continuous)
## To visually examine the correlation structure of the data, the corrplot package
## contains an excellent function of the same name.
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
corrplot(correlations, order = "hclust")
```



```
# picking out categorical predictors

categorical_vars <- c("month", "waterfront", "view", "condition", "grade", "zipcode")

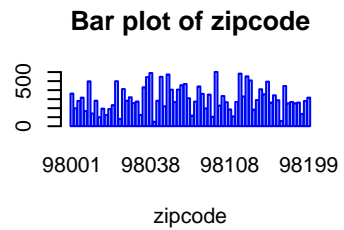
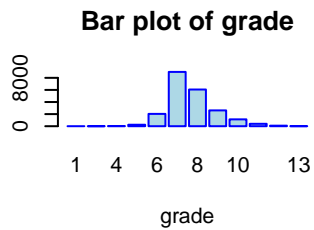
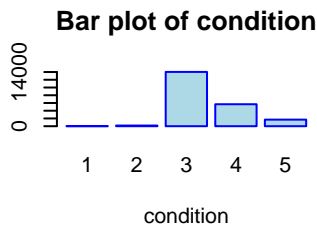
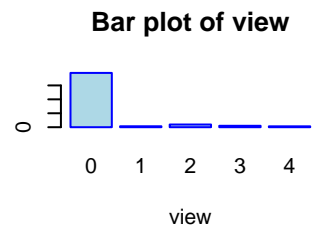
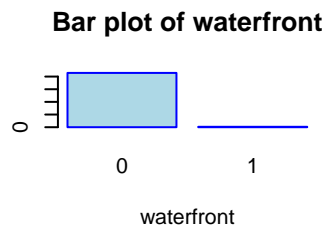
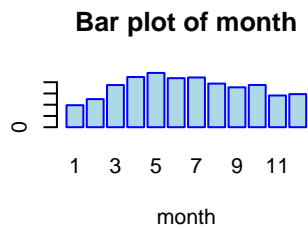
df_categorical <- hp_data_2[, categorical_vars]
```

```
#Bar plot

# Function to generate histograms for continuous predictors
create_barplot <- function(data, predictor_names) {
  # Create a layout for the bar plots
  par(mfrow = c(3, 3)) # Change the rows and columns as needed

  # Loop through each predictor and create a bar plot
  for (predictor in predictor_names) {
    barplot(table(data[[predictor]]),
            main = paste("Bar plot of", predictor),
            xlab = predictor,
            col = "lightblue",
            border = "blue")
  }
}

# Function to create histograms for continuous predictors
create_barplot(df_categorical, categorical_vars)
```



```
#Skewness Calculation
#install.packages("e1071")
library(e1071)

#skew_val <- skewness(hp_data$sqft_lot15)
#print(skew_val)

skew1 <- apply(hp_data_2, 2, skewness)
skew1
```

```
##      price      bedrooms      bathrooms      sqft_living      sqft_lot
##  4.02115735  1.97402550  0.51103663  1.47135117  13.05820621
##      floors      waterfront      view      condition      grade
##  0.61609120  11.38352768  3.39527826  1.03266128  0.77099617
##      sqft_above sqft_basement      yr_built yr_renovated      zipcode
##  1.44646367  1.57774603 -0.46974019  4.54886189  0.40560490
##      lat      long sqft_living15 sqft_lot15      year
## -0.48520312  0.88493014  1.10802746  9.50542370  0.75719402
##      month
##  0.06312141
```

```
#correlation
#install.packages("ggplot2")
library(ggplot2)
```

```
# Compute the correlation matrix
```

```
correlation_matrix <- cor(df_continuous)
```

```
# Heatmap of the correlation matrix
```

```
library(ggplot2)
```

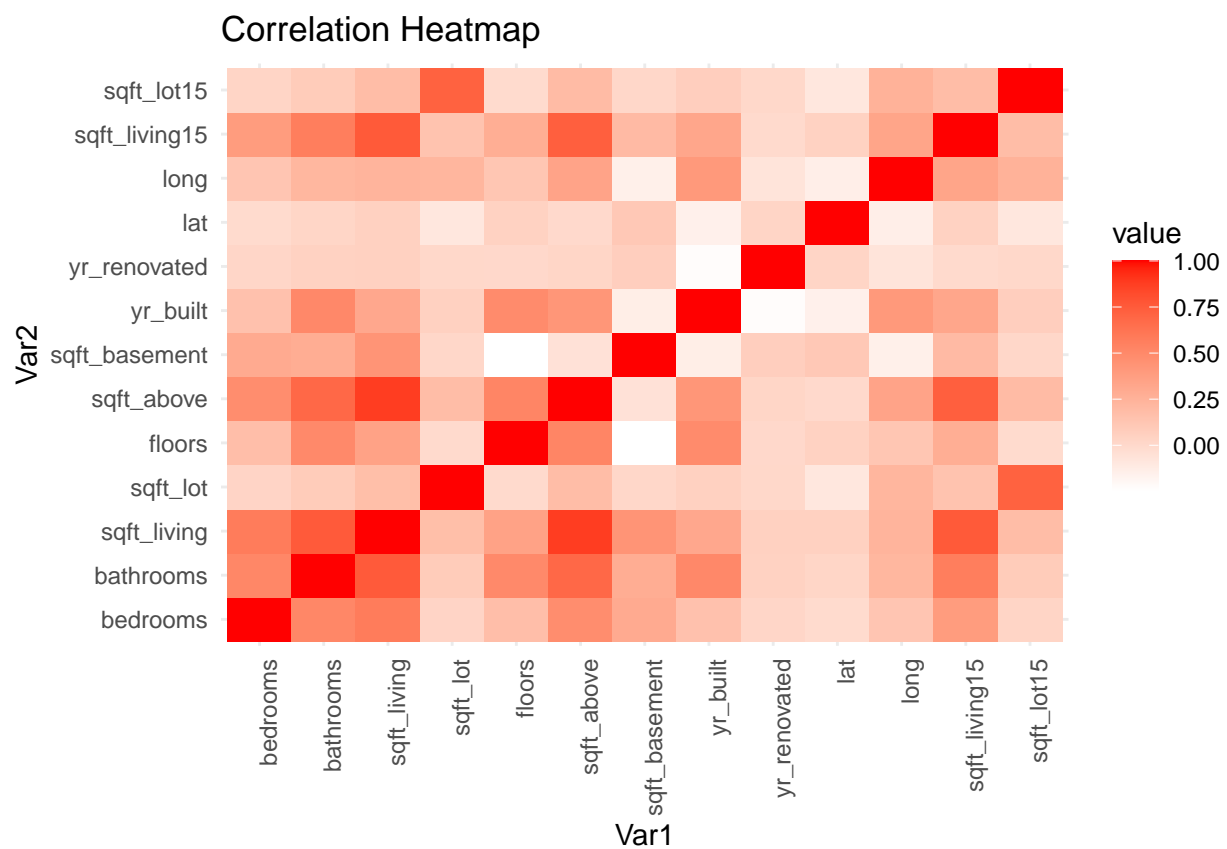
```
library(reshape2) # For melt function
```

```
# Melt the correlation matrix to a long format
```

```
correlation_data <- melt(correlation_matrix)
```

```
# heatmap
```

```
ggplot(data = correlation_data, aes(x = Var1, y = Var2, fill = value)) +  
  geom_tile() +  
  scale_fill_gradient(low = "white", high = "red") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Correlation Heatmap")
```



```
library(caret)
```

```
## Loading required package: lattice
```

```
# Look for degenerate columns:
nearZero1 <- nearZeroVar(df_continuous)
con_nearZero <- length(nearZero1)

con_nearZero
```

```
## [1] 2
```

```
new_continuous <- df_continuous[,-nearZero1]
dim(new_continuous)
```

```
## [1] 21613    11
```

```
#Near Zero for categorical predictors
nearZero2 <- nearZeroVar(df_categorical)
cat_nearZero <- length(nearZero2)

cat_nearZero
```

```
## [1] 2
```

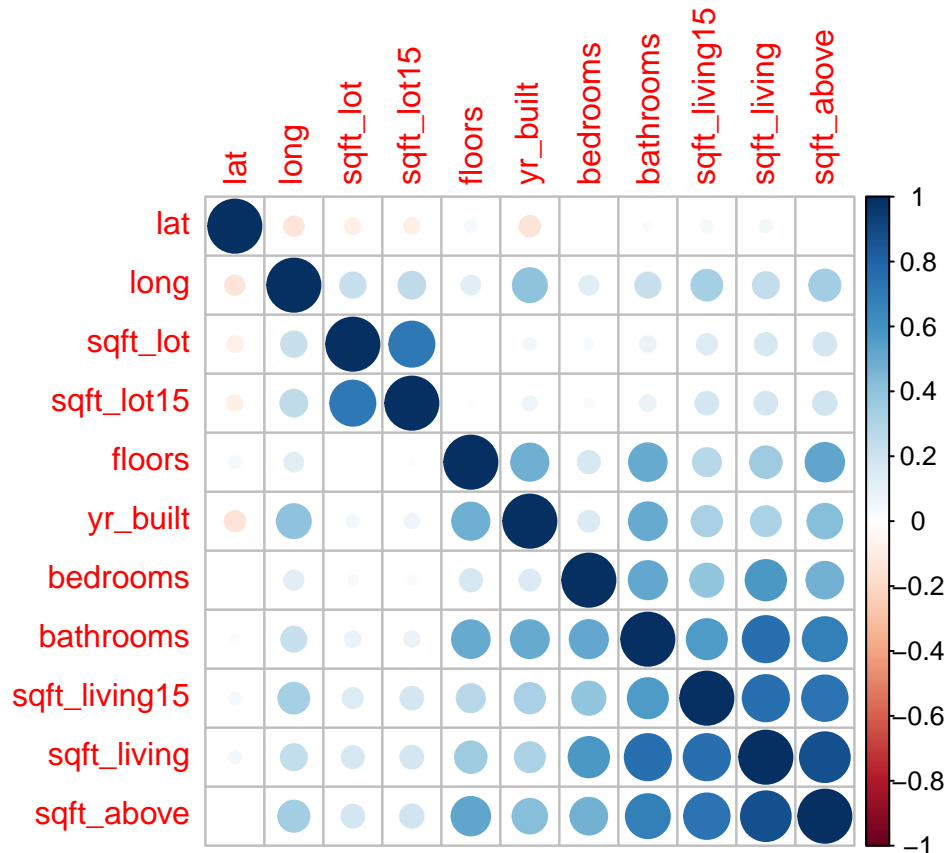
```
new_categorical <- df_categorical[,-nearZero2]
dim(new_categorical)
```

```
## [1] 21613     4
```

```
str(new_categorical)
```

```
## 'data.frame':    21613 obs. of  4 variables:
## $ month      : num  10 12 2 12 2 5 6 1 4 3 ...
## $ condition: num   3 3 3 5 3 3 3 3 3 3 ...
## $ grade      : num   7 7 6 7 8 11 7 7 7 7 ...
## $ zipcode    : num  98178 98125 98028 98136 98074 ...
```

```
# Look for strong correlations among the predictors:
library(corrplot)
corrplot(cor(new_continuous), order="hclust")
```



```
# Find which predictors we can eliminate since they have correlations that are
#"too large":
```

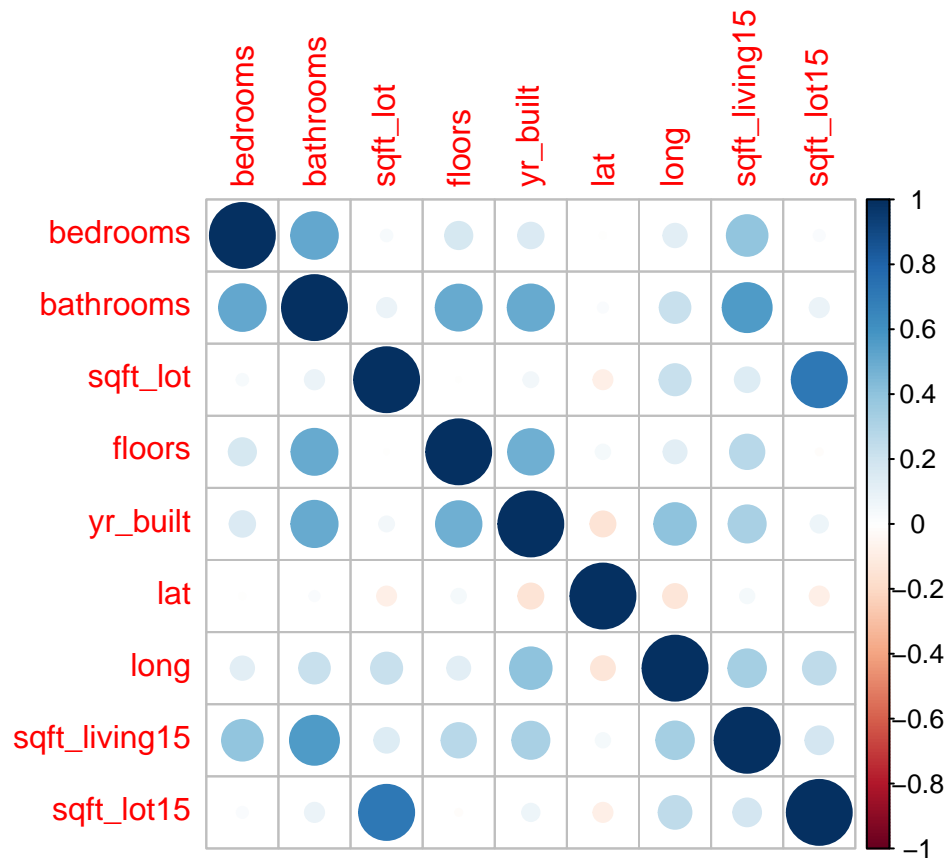
```
highCorr = findCorrelation( cor( new_continuous ), cutoff=0.75 )
#highCorr = findCorrelation( cor( new_bio ), cutoff=0.9 )
df_continuous_independent = new_continuous[,-highCorr]
#df_continuous_independent
dim(df_continuous_independent)
```

```
## [1] 21613      9
```

```
str(df_continuous_independent)
```

```
## 'data.frame':  21613 obs. of  9 variables:
## $ bedrooms   : num  3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms  : num  1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_lot    : num  5650 7242 10000 5000 8080 ...
## $ floors     : num  1 2 1 1 1 1 2 1 1 2 ...
## $ yr_built   : num  1955 1951 1933 1965 1987 ...
## $ lat        : num  47.5 47.7 47.7 47.5 47.6 ...
## $ long       : num  -122 -122 -122 -122 -122 ...
## $ sqft_living15: num  1340 1690 2720 1360 1800 ...
## $ sqft_lot15  : num  5650 7639 8062 5000 7503 ...
```

```
corrplot( cor(df_continuous_independent) ) # notice that this matrix has no values >
```



```
#cutoff=0.75 above
```

```
# Transformation
#handling the column with negative values
# Find the minimum value in the long
min_val_long <- min(df_continuous_independent$long)
print(min_val_long)
```

```
## [1] -122.519
```

```
#Adding the constant
if (min_val_long < 0) {
  df_continuous_independent$long <- df_continuous_independent$long + abs(min_val_long) + 1
}

head(df_continuous_independent)
```

```
## bedrooms bathrooms sqft_lot floors yr_built lat long sqft_living15
## 1 3 1.00 5650 1 1955 47.5112 1.262 1340
## 2 3 2.25 7242 2 1951 47.7210 1.200 1690
## 3 2 1.00 10000 1 1933 47.7379 1.286 2720
```



```
## 4      4      3.00      5000      1      1965 47.5208 1.126      1360
## 5      3      2.00      8080      1      1987 47.6168 1.474      1800
## 6      4      4.50     101930      1      2001 47.6561 1.514      4760
## sqft_lot15
## 1      5650
## 2      7639
## 3      8062
## 4      5000
## 5      7503
## 6     101930
```

```
#df_continuous_independent
```

```
#Boxcox transformation
#install.packages("car")
library(car)
```

```
## Loading required package: carData
```

```
#install.packages("MASS")
transformed_data <- df_continuous_independent
library(MASS)
```

```
#####PCA#####
trans <- preProcess(df_continuous_independent, method = c("BoxCox", "center", "scale")) ## need {caret}
#trans
```

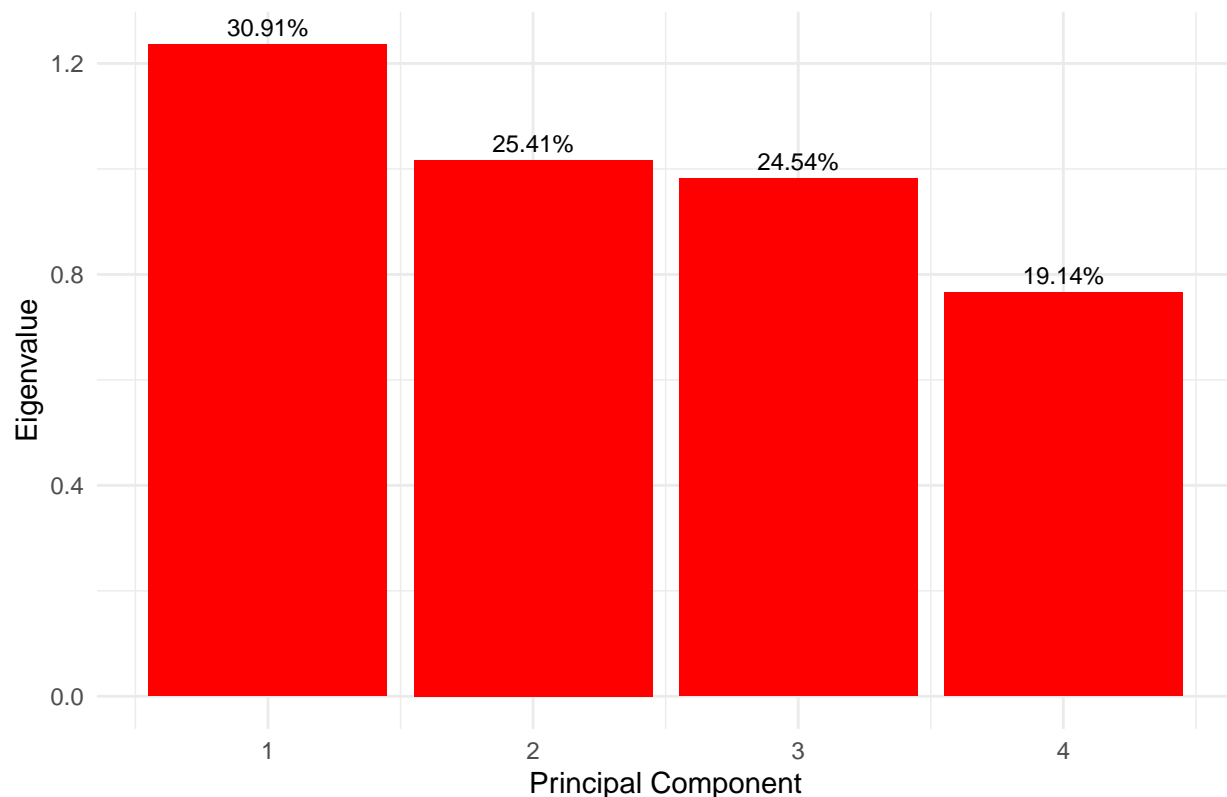
```
#####
# Load necessary library
library(ggplot2)
# Perform PCA
pca_result <- prcomp(new_categorical, scale = TRUE)

# Extract eigenvalues and calculate the percentage variance explained
eigenvalues <- pca_result$sdev^2
total_variance <- sum(eigenvalues)
percentage_variance <- eigenvalues / total_variance * 100

# Scree plot
scree_plot <- ggplot(data = data.frame(PC = 1:length(eigenvalues),
                                       Eigenvalue = eigenvalues,
                                       PercentageVariance = percentage_variance),
                    aes(x = PC, y = Eigenvalue)) +
  geom_bar(stat = "identity", fill = "red") +
  geom_text(aes(label = sprintf("%.2f%%", PercentageVariance)),
            vjust = -0.5, size = 3, color = "black") +
  labs(x = "Principal Component", y = "Eigenvalue",
       title = "Scree Plot with Percentage Variance Explained") +
  theme_minimal()

print(scree_plot)
```

Scree Plot with Percentage Variance Explained



```
#####

#### use preProcess
library(caret)
lm <- preProcess(df_continuous_independent,method=c("BoxCox")) ## need {caret} package
## Apply imputation
df_continuous_box <- predict(lm,df_continuous_independent)
#df_continuous_box
###boxcox hist.

predictor_names <- names(df_continuous_box)
#####
# Function to generate histograms with skewness for continuous predictors
create_histograms_with_skewness <- function(data, predictor_names) {
  # Create a layout for the histograms
  par(mfrow = c(2, 3)) # Change the rows and columns as needed

  # Loop through each predictor and create a histogram
  for (predictor in predictor_names) {
    # Calculate skewness
    skew <- skewness(data[[predictor]])

    # Create the histogram
    hist(data[[predictor]],
          main = paste("Histogram of", predictor, "\nSkewness:", round(skew, 2)),
          xlab = predictor,
```

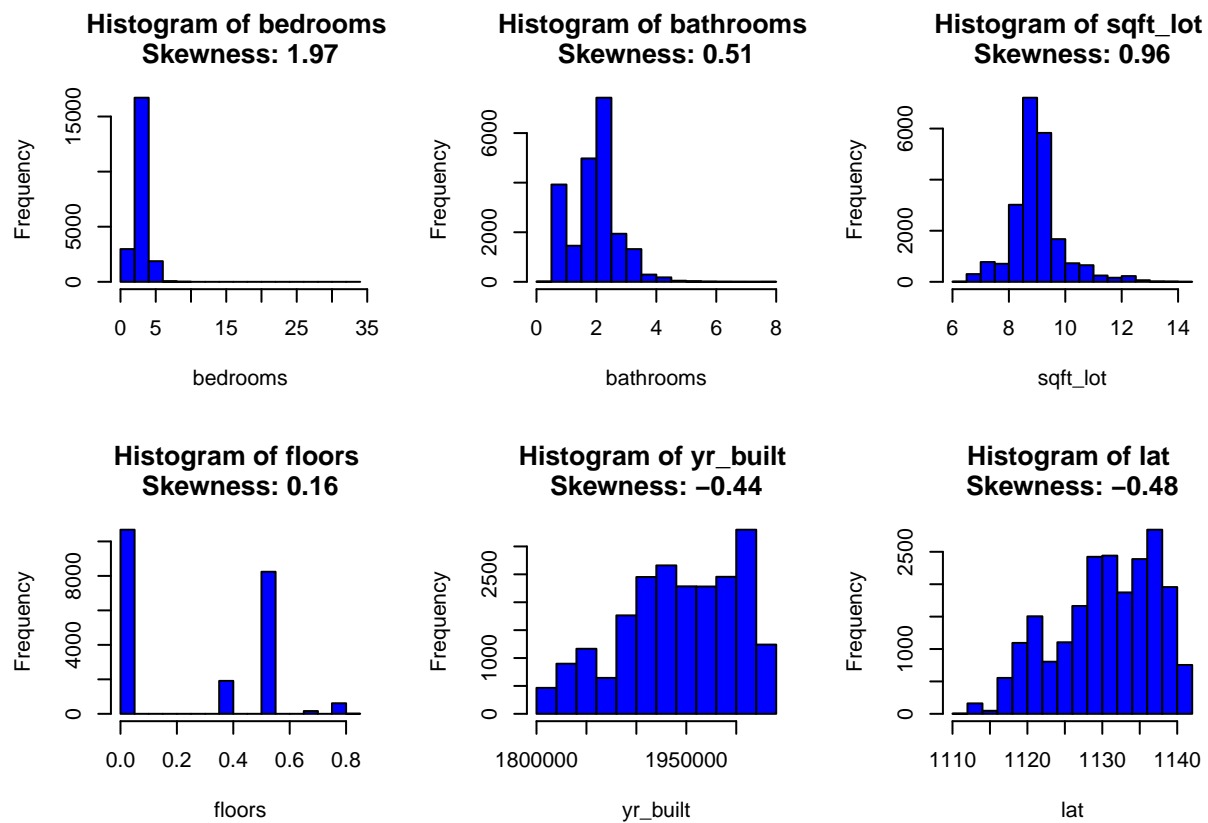
```

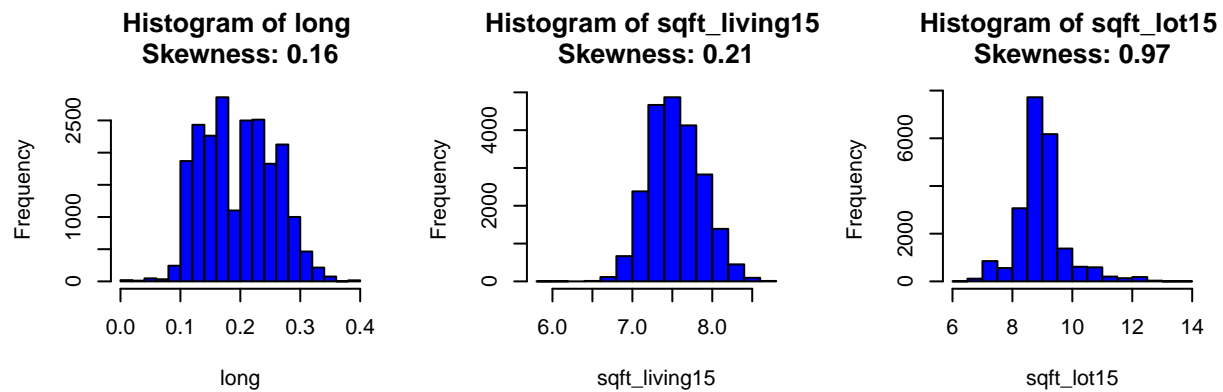
    col = "blue",
    border = "black")
  }
}

# Function to calculate skewness
skewness <- function(x) {
  n <- length(x)
  mean_val <- mean(x)
  sd_val <- sd(x)
  skewness_val <- (sum((x - mean_val)^3) / n) / (sd_val^3)
  return(skewness_val)
}

# Usage example: create histograms with skewness
create_histograms_with_skewness(df_continuous_box, predictor_names)

```





```
#####
#pca_result_extracted <- pca_result$x[, 1:4]
#pca_result_extracted
#####3
#Spatial sign
#install.packages("ICSNP")
library(ICSNP)
```

```
## Loading required package: mvtnorm
```

```
## Loading required package: ICS
```

```
# Extracting the data matrix
data_matrix <- as.matrix(df_continuous_box)

# Apply spatial sign transformation
transformed_data <- spatialSign(data_matrix)

df_transformed_1 <- as.data.frame(transformed_data)

# Boxplot after transformation

# Names of continuous predictors
predictor_names <- names(df_transformed_1)
```

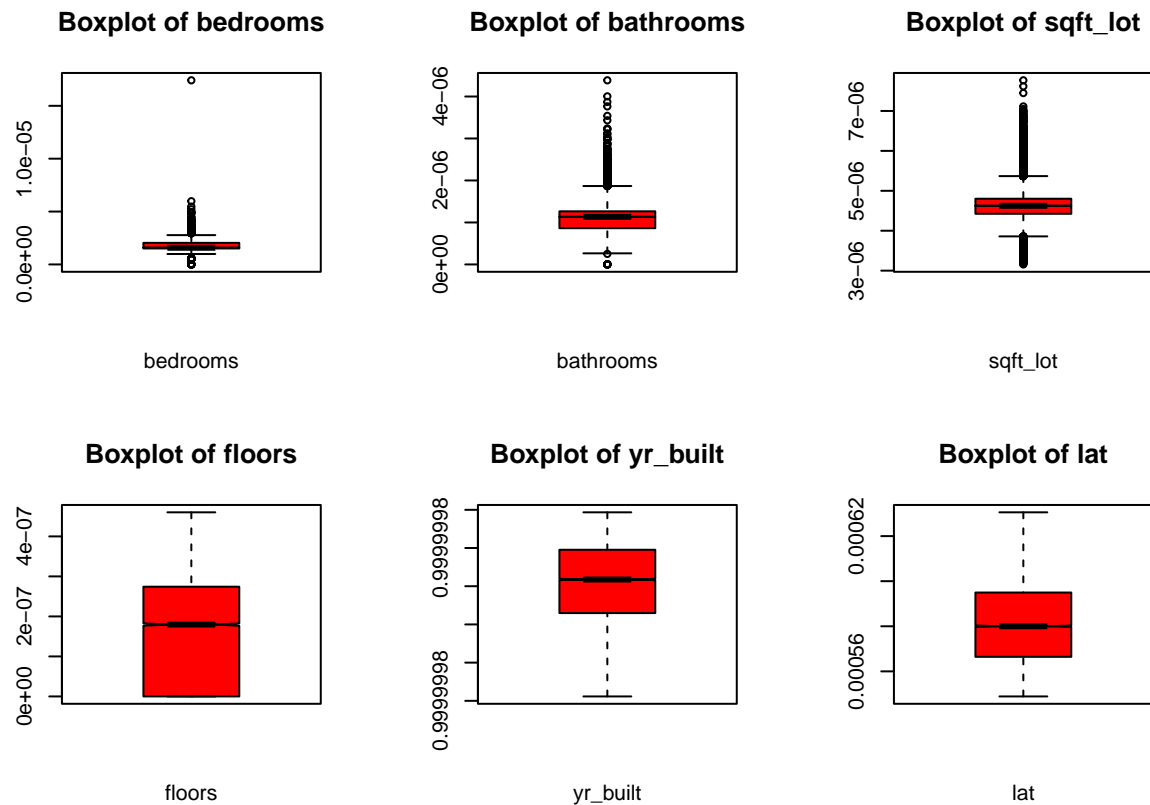
```

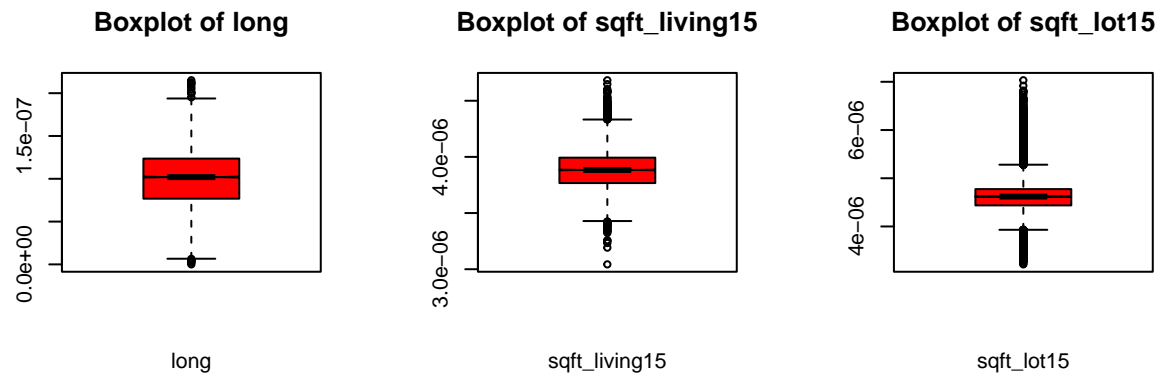
# Function to generate boxplots for all continuous predictors
create_boxplots <- function(data, predictor_names) {
  # Create a layout for the boxplots
  par(mfrow = c(2, 3)) # Change the rows and columns as needed

  # Loop through each predictor and create a boxplot
  for (predictor in predictor_names) {
    boxplot(data[[predictor]],
            main = paste("Boxplot of", predictor),
            xlab = predictor,
            col = "red",
            border = "black",
            notch = TRUE)
  }
}

# Function to create boxplots for all predictors
create_boxplots(df_transformed_1, predictor_names)

```





```
#####
#nrow(pca_result_extracted)
#ncol(pca_result_extracted)

### Combining the predictors####
combined_data <- cbind(df_continuous_box, new_categorical)
dim(combined_data)
```

```
## [1] 21613    13
```

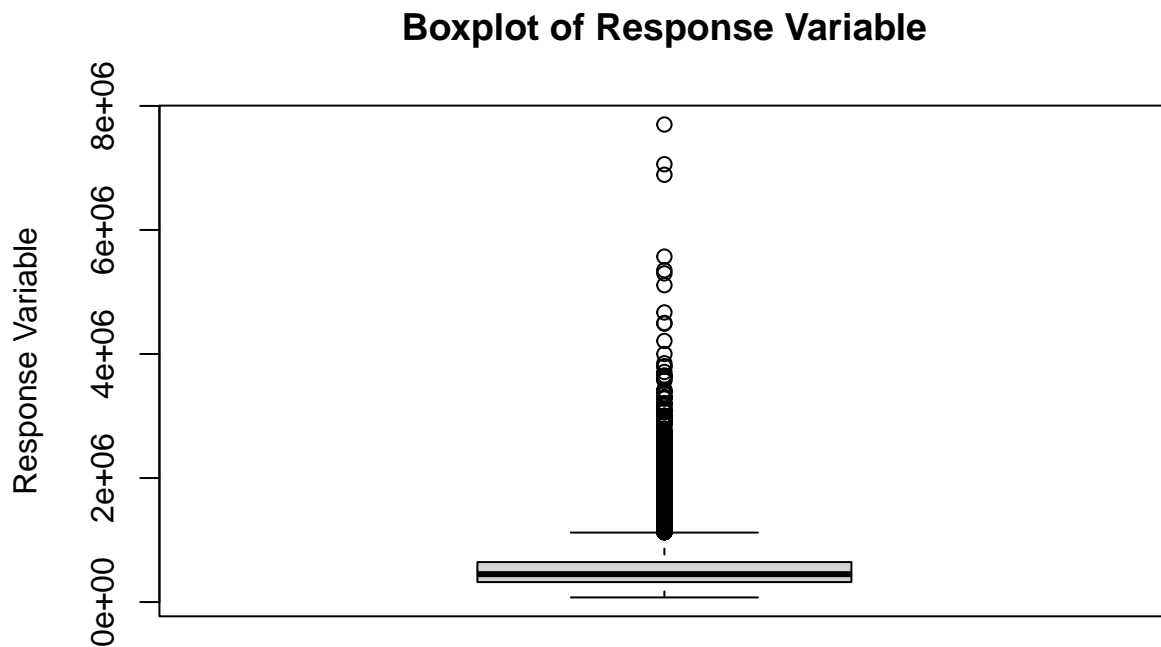
```
str(combined_data)
```

```
## 'data.frame':    21613 obs. of  13 variables:
## $ bedrooms      : num  3 3 2 4 3 4 3 3 3 3 ...
## $ bathrooms     : num  1 2.25 1 3 2 4.5 2.25 1.5 1 2.5 ...
## $ sqft_lot       : num  8.64 8.89 9.21 8.52 9 ...
## $ floors        : num  0 0.549 0 0 0 ...
## $ yr_built       : num  1911012 1903200 1868244 1930612 1974084 ...
## $ lat           : num  1128 1138 1139 1129 1133 ...
## $ long          : num  0.186 0.153 0.198 0.106 0.27 ...
## $ sqft_living15 : num  7.2 7.43 7.91 7.22 7.5 ...
## $ sqft_lot15    : num  8.64 8.94 8.99 8.52 8.92 ...
## $ month         : num  10 12 2 12 2 5 6 1 4 3 ...
## $ condition     : num  3 3 3 5 3 3 3 3 3 3 ...
```

```
## $ grade      : num  7 7 6 7 8 11 7 7 7 7 ...
## $ zipcode    : num  98178 98125 98028 98136 98074 ...
```

```
####Extracting the response variable from the data####
price <- hp_data$price
#price
```

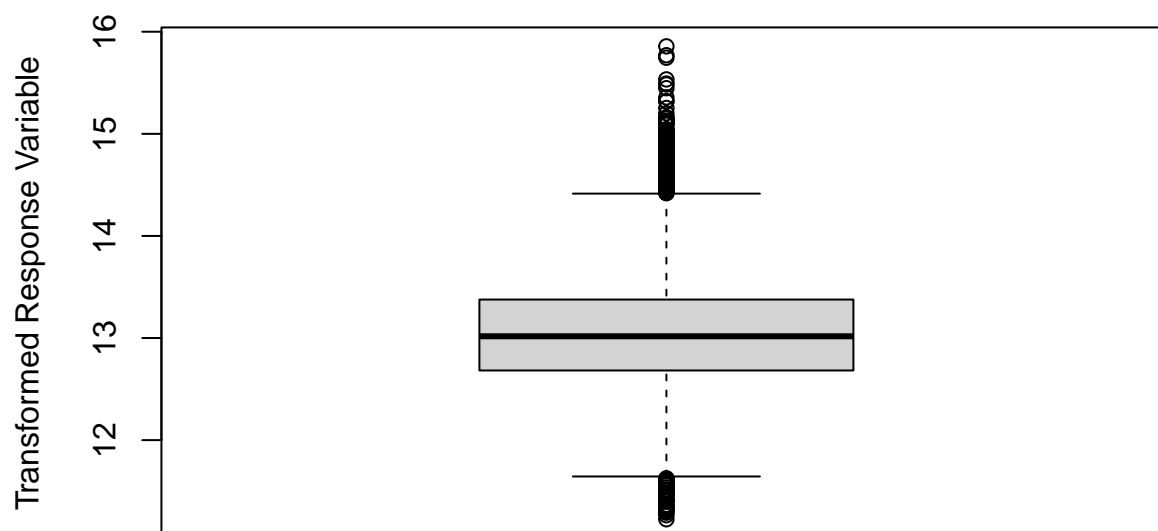
```
# Create a boxplot for the response variable
boxplot(price, main="Boxplot of Response Variable", ylab="Response Variable")
```



```
#####Tranforming price#####
transformed_price <- log1p(price)
```

```
#####Boxplot for Price after transformation#####
boxplot(transformed_price, main="Boxplot of Log-transformed Response Variable", ylab="Transformed Response Variable")
```

## Boxplot of Log-transformed Response Variable



```
price <- transformed_price
```

```
####Data Splitting###
set.seed(100)
data_split <- createDataPartition(price, p = 0.8, list = FALSE)
hp_train <- combined_data[data_split,]
price_train <- price[data_split]
hp_test <- combined_data[-data_split,]
price_test <- price[-data_split]
####to reduce computation time, i would use 700 samples
hp_train <- hp_train[1:700,]
price_train <- price_train[1:700]
hp_test <- hp_test[1:700,]
price_test <- price_test[1:700]
control <- trainControl(method = "repeatedcv", repeats = 5)
```

```
#####Linear Models####
###ordinary linear model
lmModel <- train(hp_train, price_train, method = "lm", trControl = control, length=5)
```

```
## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded
```

```
## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded
```







```

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'length' will be disregarded

```

```
lmModel
```

```

## Linear Regression
##
## 700 samples
## 13 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 631, 629, 630, 629, 630, 631, ...
## Resampling results:
##
## RMSE      Rsquared  MAE
## 0.2718258  0.73549   0.2038356

```

```
##
## Tuning parameter 'intercept' was held constant at a value of TRUE

linear_pred <- predict(lmModel , hp_test)

postResample(linear_pred, price_test)

##      RMSE  Rsquared      MAE
## 0.2808385 0.7266048 0.2143352

#####training
linear_pred2 <- predict(lmModel , hp_train)

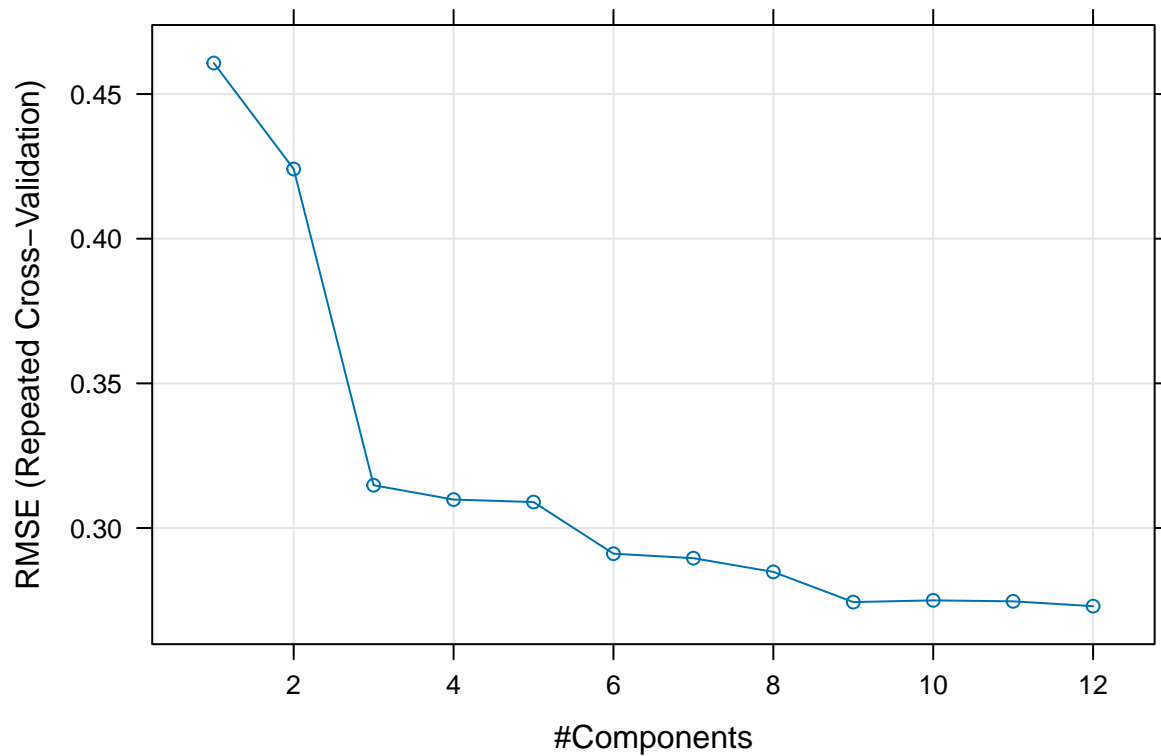
postResample(linear_pred2, price_train)

##      RMSE  Rsquared      MAE
## 0.2673853 0.7422630 0.1996422

####PCR####
pcr <- train(hp_train, price_train, method = "pcr",
             preProcess = c("center", "scale"),
             trControl = control(),
             #tuneGrid = data.frame(ncomp = num_components),
             tuneLength = 40)
pcr

## Principal Component Analysis
##
## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 631, 630, 629, 630, 629, 630, ...
## Resampling results across tuning parameters:
##
##  ncomp  RMSE      Rsquared  MAE
##    1     0.4607254 0.2366007 0.3759502
##    2     0.4240923 0.3537026 0.3409292
##    3     0.3147865 0.6439748 0.2458457
##    4     0.3098273 0.6548860 0.2398265
##    5     0.3089654 0.6574231 0.2388155
##    6     0.2911106 0.6964794 0.2184721
##    7     0.2895646 0.6997303 0.2168537
##    8     0.2848228 0.7097426 0.2151369
##    9     0.2743858 0.7299178 0.2058997
##   10     0.2749771 0.7288114 0.2066388
##   11     0.2746486 0.7293336 0.2063519
##   12     0.2729808 0.7329419 0.2042022
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 12.
```

```
plot(pcr)
```



```
#predicting on testing
predict_pcr <- predict(pcr, hp_test)
postResample(predict_pcr, price_test)
```

```
##      RMSE  Rsquared    MAE
## 0.2767417 0.7354336 0.2112970
```

```
#####predict on training
predict_pcr2 <- predict(pcr, hp_train)
postResample(predict_pcr2, price_train)
```

```
##      RMSE  Rsquared    MAE
## 0.2685348 0.7400422 0.1999526
```

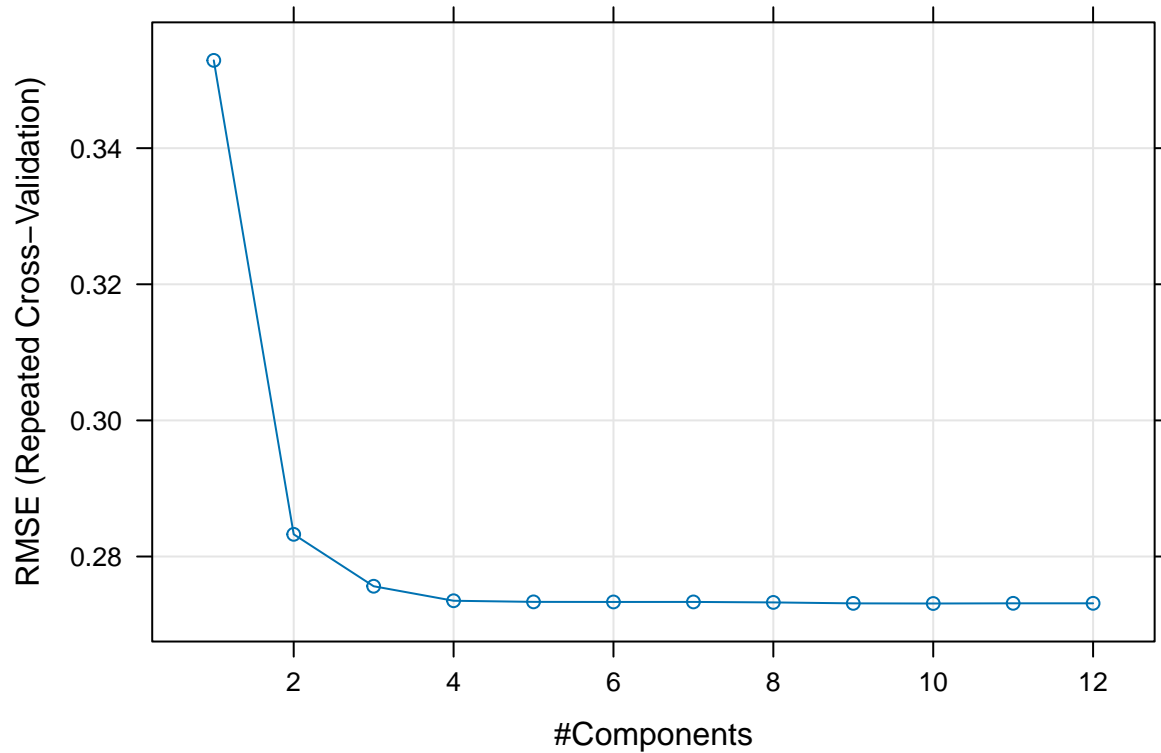
```
###PLS###
#PLS MODEL
model_pls <- train(hp_train, price_train, method = "pls",
  tuneLength = 40,
  preProcess = c("center", "scale"),
  trControl = trainControl(method = "repeatedcv", repeats = 5))
# train_()_function_output_and_tuning_parameter_plot
print(model_pls)
```

```

## Partial Least Squares
##
## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 630, 631, 630, 630, 630, 629, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      0.3529009  0.5557225  0.2820066
##   2      0.2832611  0.7159497  0.2171435
##   3      0.2756318  0.7309569  0.2085879
##   4      0.2734987  0.7349481  0.2037015
##   5      0.2733335  0.7352305  0.2049855
##   6      0.2733241  0.7353410  0.2050019
##   7      0.2733287  0.7353637  0.2047030
##   8      0.2732544  0.7355084  0.2048047
##   9      0.2731133  0.7357948  0.2048471
##  10      0.2730960  0.7358373  0.2047509
##  11      0.2731237  0.7357819  0.2047780
##  12      0.2731233  0.7357821  0.2047858
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 10.

```

```
plot(model_pls)
```



```
#predicting on testing
predictions_pls <- predict(model_pls, hp_test)
postResample(predictions_pls, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.2807923 0.7267022 0.2143009
```

```
#####predicting on training predictions_pls2 <- predict(model_pls, hp_train) postResam-
ple(predictions_pls2, price_train)
```

```
#####variable importance
varImp(model_pls)
```

```
##
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:caret':
##
##      R2
```

```
## The following object is masked from 'package:corrplot':
##
##      corrplot
```

```
## The following object is masked from 'package:stats':
##
##   loadings
```

```
## pls variable importance
```

```
##
##           Overall
## grade          100.00
## sqft_living15   89.81
## lat             85.94
## bathrooms       70.94
## bedrooms        43.71
## floors          40.55
## sqft_lot        26.34
## sqft_lot15      26.01
## yr_built        17.88
## long            16.75
## zipcode         14.13
## condition       12.61
## month           0.00
```

```
#####
ridgeGrid <- data.frame(.lambda = seq(0.05, .9, length = 15))
lassoGrid <- expand.grid(alpha = 1, lambda = c(0, 0.01, 0.4))
enetGrid <- expand.grid(.lambda = c(0, 0.01, .1), .fraction = seq(.05, 1, length = 20))
#####
```

```
#####Ridge###
ridgeModel <- train(hp_train, price_train,
  method = "ridge",
  tuneGrid = ridgeGrid,
  trControl = control,
  preProc = c("center", "scale"))
ridgeModel
```

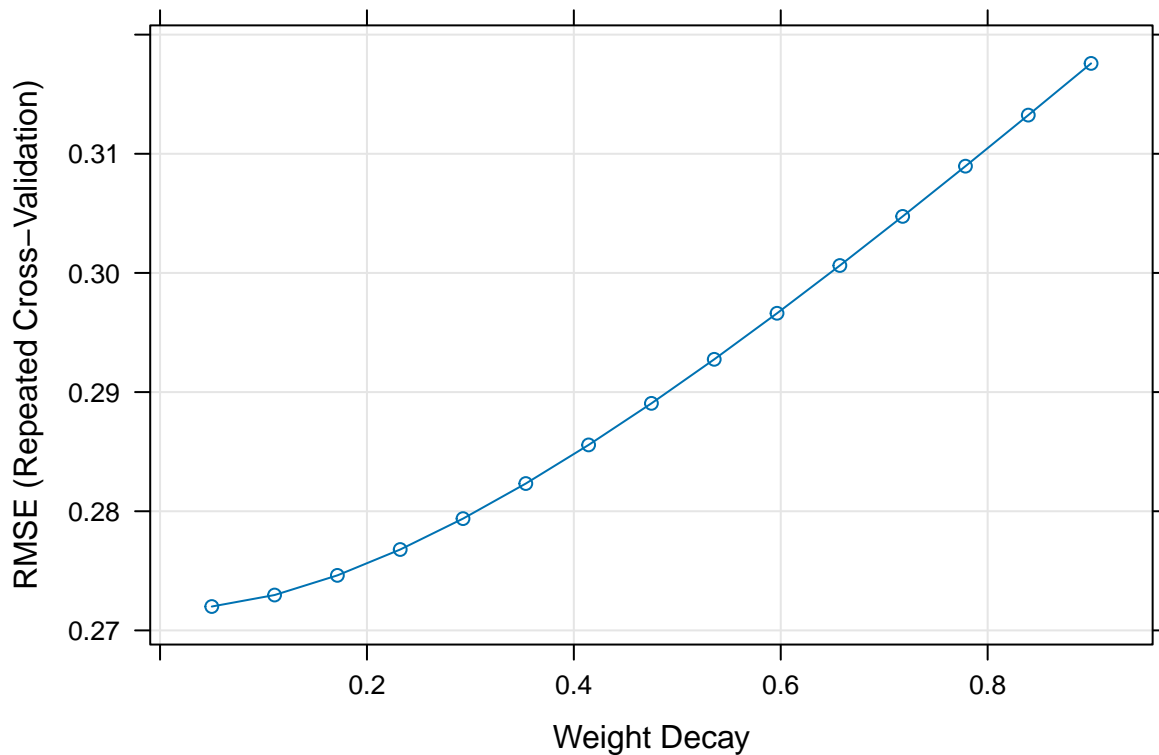
```
## Ridge Regression
```

```
##
## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 630, 630, 631, 631, 629, 630, ...
## Resampling results across tuning parameters:
##
##   lambda      RMSE      Rsquared    MAE
## 0.0500000  0.2720028  0.7356791  0.2045485
## 0.1107143  0.2729663  0.7349995  0.2058899
## 0.1714286  0.2746165  0.7337256  0.2077320
## 0.2321429  0.2767866  0.7321375  0.2100386
## 0.2928571  0.2793800  0.7303624  0.2125782
## 0.3535714  0.2823256  0.7284716  0.2153372
## 0.4142857  0.2855663  0.7265093  0.2182088
```



```
## 0.4750000 0.2890544 0.7245048 0.2212403
## 0.5357143 0.2927491 0.7224780 0.2244675
## 0.5964286 0.2966153 0.7204432 0.2277743
## 0.6571429 0.3006228 0.7184106 0.2311255
## 0.7178571 0.3047451 0.7163878 0.2345153
## 0.7785714 0.3089594 0.7143806 0.2379512
## 0.8392857 0.3132458 0.7123932 0.2414867
## 0.9000000 0.3175869 0.7104289 0.2450617
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was lambda = 0.05.
```

```
plot(ridgeModel)
```



```
#predicting on testing
predictions_ridge <- predict(ridgeModel, hp_test)
postResample(predictions_ridge, price_test)
```

```
## RMSE Rsquared MAE
## 0.2786152 0.7302364 0.2129907
```

```
#####predicting on training
predictions_ridge2 <- predict(ridgeModel, hp_train)
postResample(predictions_ridge2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.2678994 0.7414924 0.2005945
```

```
###LASSO###
```

```
library(elasticnet)
```

```
## Loading required package: lars
```

```
## Loaded lars 1.3
```

```
library(caret)
lasso_Model <- train(hp_train, price_train, method = "lasso",
                    trControl = control,
                    )
lasso_Model
```

```
## The lasso
```

```
##
```

```
## 700 samples
```

```
## 13 predictor
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold, repeated 5 times)
```

```
## Summary of sample sizes: 631, 629, 629, 630, 630, 631, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
## fraction RMSE      Rsquared  MAE
```

```
## 0.1      0.4602384 0.4814317 0.3661942
```

```
## 0.5      0.3019407 0.7027341 0.2276939
```

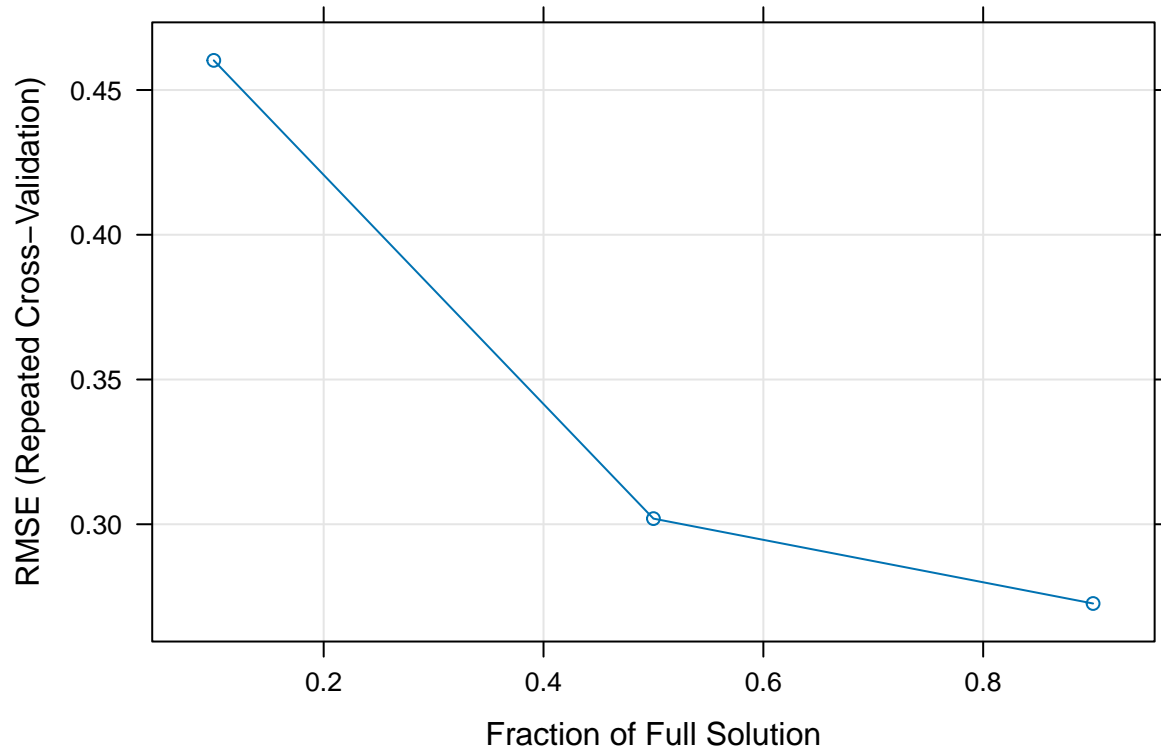
```
## 0.9      0.2726201 0.7365682 0.2043247
```

```
##
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final value used for the model was fraction = 0.9.
```

```
plot(lasso_Model)
```



```
#predicting on testing
predictions_lasso <- predict(lasso_Model, hp_test)
postResample(predictions_lasso, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.2792168 0.7310327 0.2127976
```

```
#####predicting on training
predictions_lasso2 <- predict(lasso_Model, hp_train)
postResample(predictions_lasso2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.2678241 0.7414997 0.1998311
```

```
###Elastic Net###
enetGrid <- expand.grid(.fraction = seq(0, 1, by = 0.1),
#                      .lambda = seq(0.001, 0.1, length = 15))
set.seed(100)
elastic <- train(hp_train, price_train, method = "enet",
                tuneGrid = enetGrid, trControl = control, preProc = c("center", "scale"))
elastic
```

```
## Elasticnet
##
```

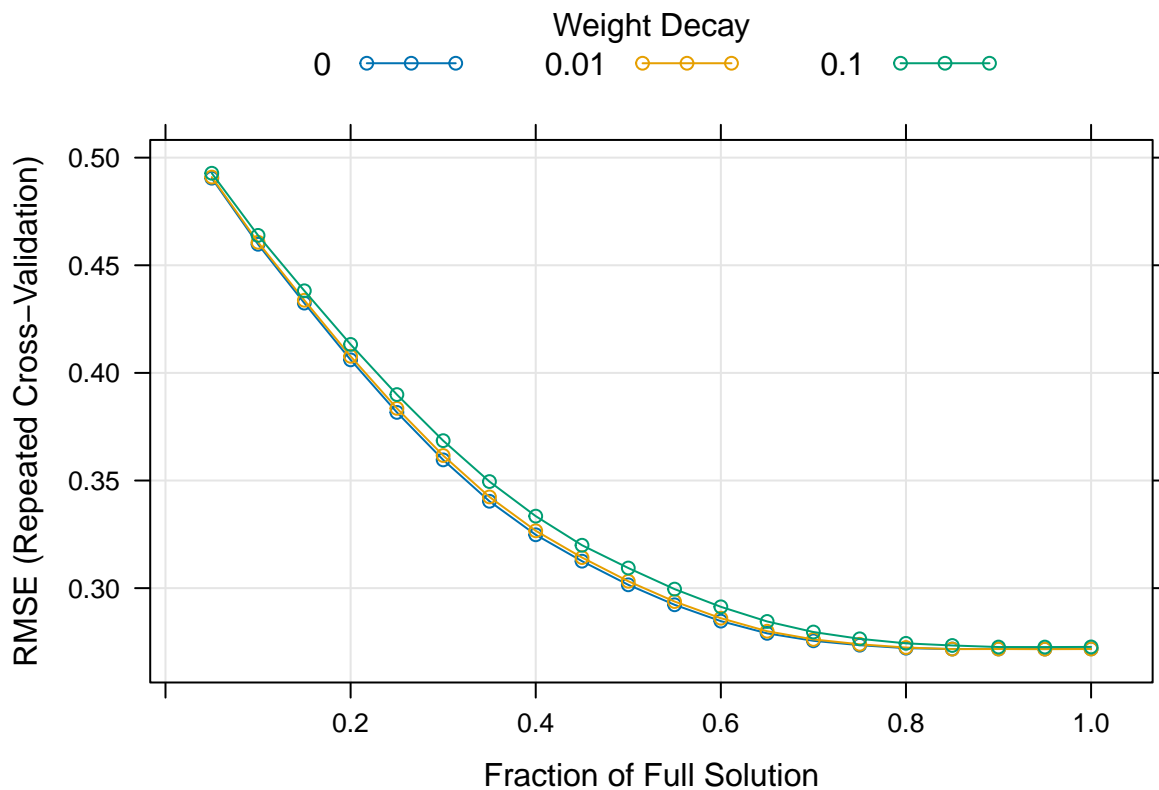
```

## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 630, 630, 630, 629, 630, 631, ...
## Resampling results across tuning parameters:
##
##   lambda  fraction  RMSE      Rsquared  MAE
##   0.00    0.05      0.4904448  0.4754957  0.3891687
##   0.00    0.10      0.4597429  0.4821646  0.3659046
##   0.00    0.15      0.4324210  0.5576490  0.3429588
##   0.00    0.20      0.4060578  0.6152517  0.3192999
##   0.00    0.25      0.3816458  0.6443481  0.2976879
##   0.00    0.30      0.3595866  0.6596708  0.2781377
##   0.00    0.35      0.3403577  0.6680054  0.2613268
##   0.00    0.40      0.3247875  0.6728969  0.2478858
##   0.00    0.45      0.3125231  0.6848091  0.2369166
##   0.00    0.50      0.3015461  0.7019037  0.2273497
##   0.00    0.55      0.2922651  0.7142923  0.2189990
##   0.00    0.60      0.2846726  0.7232194  0.2124413
##   0.00    0.65      0.2790072  0.7290536  0.2078069
##   0.00    0.70      0.2755277  0.7320386  0.2053108
##   0.00    0.75      0.2734994  0.7341570  0.2043086
##   0.00    0.80      0.2721177  0.7358689  0.2038603
##   0.00    0.85      0.2717174  0.7361716  0.2039185
##   0.00    0.90      0.2717511  0.7359660  0.2040047
##   0.00    0.95      0.2717130  0.7358746  0.2040438
##   0.00    1.00      0.2718056  0.7356556  0.2042030
##   0.01    0.05      0.4909795  0.4754957  0.3895749
##   0.01    0.10      0.4606878  0.4827014  0.3666083
##   0.01    0.15      0.4337408  0.5541848  0.3441809
##   0.01    0.20      0.4077073  0.6127624  0.3208200
##   0.01    0.25      0.3835391  0.6426200  0.2993997
##   0.01    0.30      0.3616142  0.6584822  0.2799970
##   0.01    0.35      0.3423787  0.6671997  0.2630793
##   0.01    0.40      0.3266576  0.6722790  0.2496044
##   0.01    0.45      0.3142055  0.6815436  0.2384116
##   0.01    0.50      0.3032092  0.6988889  0.2288899
##   0.01    0.55      0.2937783  0.7118733  0.2203942
##   0.01    0.60      0.2860304  0.7212626  0.2136093
##   0.01    0.65      0.2799977  0.7278081  0.2086082
##   0.01    0.70      0.2761880  0.7312853  0.2057806
##   0.01    0.75      0.2739212  0.7334657  0.2045139
##   0.01    0.80      0.2724117  0.7354141  0.2039462
##   0.01    0.85      0.2716842  0.7362501  0.2038880
##   0.01    0.90      0.2717159  0.7360746  0.2040490
##   0.01    0.95      0.2716519  0.7360728  0.2040885
##   0.01    1.00      0.2717192  0.7359165  0.2042419
##   0.10    0.05      0.4927487  0.4754957  0.3909144
##   0.10    0.10      0.4638772  0.4901088  0.3690138
##   0.10    0.15      0.4381461  0.5413684  0.3484202
##   0.10    0.20      0.4132337  0.6025279  0.3261280
##   0.10    0.25      0.3899243  0.6348524  0.3053890

```

```
## 0.10 0.30 0.3685249 0.6526888 0.2865739
## 0.10 0.35 0.3495123 0.6627989 0.2696959
## 0.10 0.40 0.3334509 0.6696658 0.2559720
## 0.10 0.45 0.3199313 0.6750641 0.2441608
## 0.10 0.50 0.3093510 0.6835721 0.2347860
## 0.10 0.55 0.2995568 0.6987571 0.2262143
## 0.10 0.60 0.2913301 0.7101973 0.2186023
## 0.10 0.65 0.2845308 0.7190882 0.2126563
## 0.10 0.70 0.2796241 0.7250438 0.2087182
## 0.10 0.75 0.2764594 0.7287519 0.2064986
## 0.10 0.80 0.2743584 0.7315175 0.2053834
## 0.10 0.85 0.2733682 0.7331334 0.2052334
## 0.10 0.90 0.2726475 0.7346761 0.2053205
## 0.10 0.95 0.2726378 0.7352001 0.2057032
## 0.10 1.00 0.2726929 0.7355086 0.2059224
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were fraction = 0.95 and lambda = 0.01.
```

```
plot(elastic)
```



```
#predicting on the testing
predictions_elastic <- predict(elastic, hp_test)
postResample(predictions_elastic, price_test)
```

```
## RMSE Rsquared MAE
```

```
## 0.2793722 0.7299261 0.2131267
```

```
#####predicting on training
predictions_elastic2 <- predict(elastic, hp_train)
postResample(predictions_elastic2, price_train)
```

```
##      RMSE Rsquared      MAE
## 0.2675900 0.7418724 0.1998949
```

```
#####Nonlinear models####
```

```
#####Neural Networks
```

```
library(caret)
```

```
nnetGrid <- expand.grid(
  size = 1:10,
  decay = c(0, 0.01, 0.1)
)
```

```
set.seed(0)
```

```
# Define your control object for training (if not already defined)
# control <- trainControl(method = "cv", number = 10)
```

```
nnet <- train(
  x = hp_train,      # Predictor variables
  y = price_train,   # Target variable
  method = "nnet",
  tuneGrid = nnetGrid,
  trControl = control,
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(hp_train) + 1) + 10 + 1,
  maxit = 500
)
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
nnet
```

```
## Neural Network
```

```
##
```

```
## 700 samples
```

```
## 13 predictor
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold, repeated 5 times)
```

```
## Summary of sample sizes: 629, 630, 631, 630, 630, 630, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
##   size decay RMSE      Rsquared      MAE
```

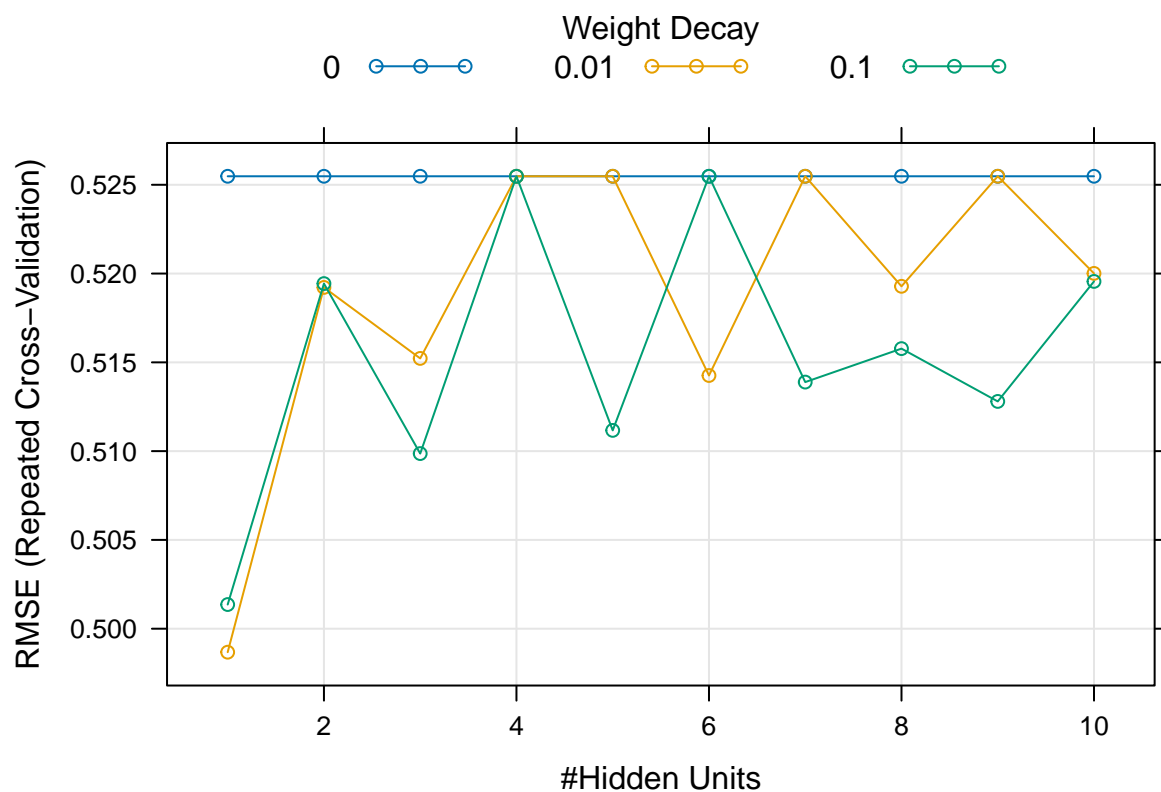
```
##      1      0.00      0.5254793      NaN      0.4149076
##      1      0.01      0.4986761      0.378897919      0.3929568
##      1      0.10      0.5013606      0.445618439      0.3938241
##      2      0.00      0.5254793      NaN      0.4149076
##      2      0.01      0.5192202      0.154838475      0.4094815
##      2      0.10      0.5194362      0.272109193      0.4093881
##      3      0.00      0.5254793      NaN      0.4149076
##      3      0.01      0.5152212      0.520310919      0.4067999
##      3      0.10      0.5098626      0.458880817      0.4008144
##      4      0.00      0.5254793      NaN      0.4149076
##      4      0.01      0.5254792      NaN      0.4149033
##      4      0.10      0.5254776      NaN      0.4148674
##      5      0.00      0.5254793      NaN      0.4149076
##      5      0.01      0.5254792      0.012048380      0.4149040
##      5      0.10      0.5111685      0.439665262      0.4031372
##      6      0.00      0.5254793      NaN      0.4149076
##      6      0.01      0.5142587      0.776930219      0.4060332
##      6      0.10      0.5254774      NaN      0.4148742
##      7      0.00      0.5254793      NaN      0.4149076
##      7      0.01      0.5254793      0.007730707      0.4149043
##      7      0.10      0.5138881      0.557160723      0.4052510
##      8      0.00      0.5254793      NaN      0.4149076
##      8      0.01      0.5192831      0.814627816      0.4102718
##      8      0.10      0.5157733      0.380394702      0.4070678
##      9      0.00      0.5254793      NaN      0.4149076
##      9      0.01      0.5254793      NaN      0.4149047
##      9      0.10      0.5127995      0.325174227      0.4049739
##     10      0.00      0.5254793      NaN      0.4149076
##     10      0.01      0.5200098      0.829732010      0.4104795
##     10      0.10      0.5195582      0.739021237      0.4099205
```

```
##
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were size = 1 and decay = 0.01.
```

```
plot(nnet)
```



```
#predicting on testing
nnet_Pred = predict(nnet, hp_test)
postResample(nnet_Pred, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.5330651      NA 0.4241900
```

```
#####predicting on training
nnet_Pred2 = predict(nnet, hp_train)
postResample(nnet_Pred2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.5266828      NA 0.4148505
```

```
####Average Neural Networks#
library(nnet)
```

```
avgnnetGrid <- expand.grid(.decay = c(0, 0.01, .1),
                          .size = c(1:10),
                          ## The next option is to use bagging (see the
                          ## next chapter) instead of different random
                          ## seeds.
                          .bag = FALSE)
```



```

avgnnnet <- train(hp_train, price_train,
  method = "avNNet",
  tuneGrid = avgnnnetGrid,
  trControl = control,
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(hp_train) + 1) + 10 + 1,
  maxit = 500
)

```

```
## Warning: executing %dopar% sequentially: no parallel backend registered
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
avgnnnet
```

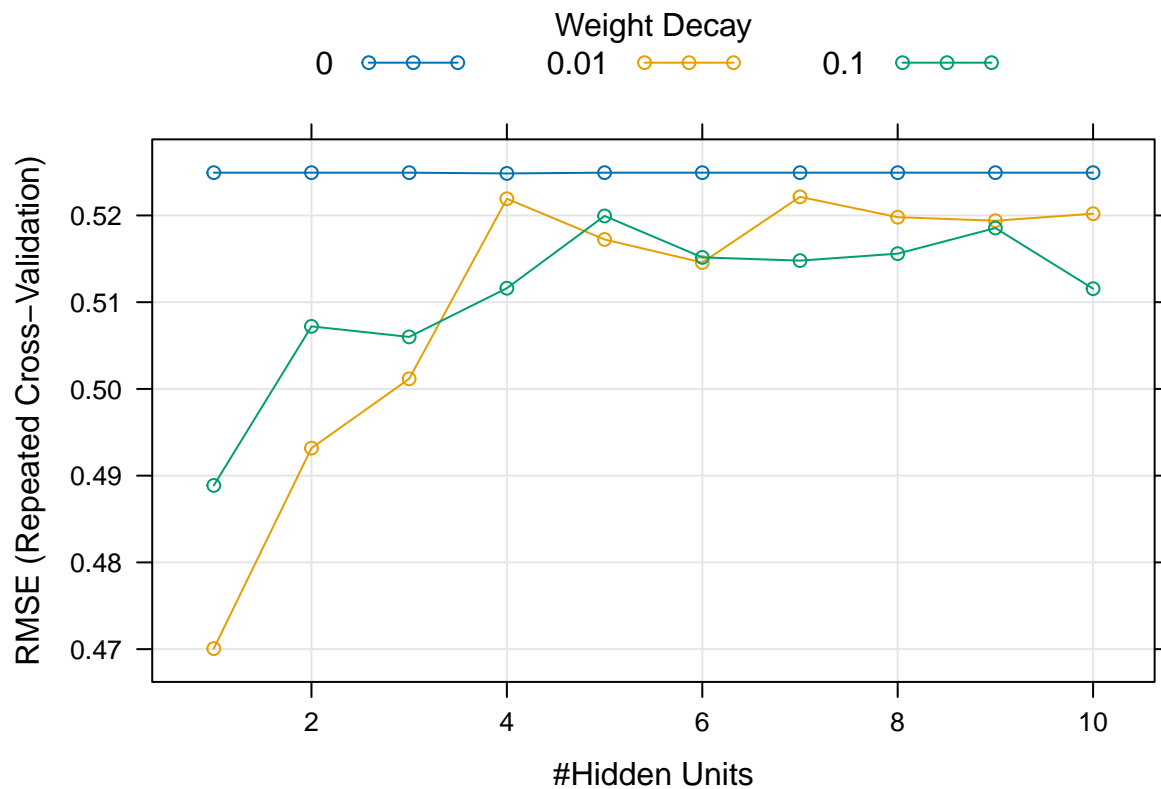
```

## Model Averaged Neural Network
##
## 700 samples
## 13 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 630, 631, 631, 630, 630, 629, ...
## Resampling results across tuning parameters:
##
##   decay  size  RMSE      Rsquared  MAE
##   0.00   1    0.5249323      NaN      0.4148105
##   0.00   2    0.5249323      NaN      0.4148105
##   0.00   3    0.5249323      NaN      0.4148105
##   0.00   4    0.5248435  0.1146402  0.4147709
##   0.00   5    0.5249323      NaN      0.4148105
##   0.00   6    0.5249323      NaN      0.4148105
##   0.00   7    0.5249323      NaN      0.4148105
##   0.00   8    0.5249323      NaN      0.4148105
##   0.00   9    0.5249323      NaN      0.4148105
##   0.00  10    0.5249323      NaN      0.4148105
##   0.01   1    0.4700634  0.4770950  0.3703375
##   0.01   2    0.4931699  0.5207533  0.3884211
##   0.01   3    0.5011640  0.4582326  0.3951692
##   0.01   4    0.5219219  0.1921429  0.4121056
##   0.01   5    0.5172320  0.3656523  0.4085232
##   0.01   6    0.5145713  0.5986591  0.4061995
##   0.01   7    0.5221439  0.2992609  0.4123563
##   0.01   8    0.5198004  0.6307389  0.4104433
##   0.01   9    0.5193919  0.6404085  0.4101466
##   0.01  10    0.5202013  0.3915647  0.4108288
##   0.10   1    0.4888753  0.5101427  0.3855241
##   0.10   2    0.5072126  0.6880527  0.4003515
##   0.10   3    0.5059972  0.5400402  0.3994097
##   0.10   4    0.5116111  0.7098194  0.4040158

```

```
## 0.10 5 0.5199410 0.4954920 0.4105773
## 0.10 6 0.5151654 0.5170275 0.4070512
## 0.10 7 0.5147928 0.4925426 0.4066995
## 0.10 8 0.5155942 0.4067898 0.4068979
## 0.10 9 0.5185507 0.4742320 0.4096459
## 0.10 10 0.5115575 0.6468320 0.4040671
##
## Tuning parameter 'bag' was held constant at a value of FALSE
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 1, decay = 0.01 and bag = FALSE.
```

```
plot(avgnnet)
```



```
#predicting on testing
avgnnet_Pred = predict(avgnnet, hp_test)
postResample(avgnnet_Pred, price_test)
```

```
##          RMSE    Rsquared      MAE
## 0.53306505 0.00235288 0.42418917
```

```
#####predicting on training
avgnnet_Pred2 = predict(avgnnet, hp_train)
postResample(avgnnet_Pred2, price_train)
```

```
##          RMSE    Rsquared      MAE
## 0.5266827812 0.0007864283 0.4148495961
```

```
###SVM###
```

```
svmTuned <- train(hp_train, price_train,  
                  method = "svmRadial",  
                  preProc = c("center", "scale"),  
                  tuneLength = 10,  
                  trControl = control)
```

```
svmTuned
```

```
## Support Vector Machines with Radial Basis Function Kernel
```

```
##
```

```
## 700 samples
```

```
## 13 predictor
```

```
##
```

```
## Pre-processing: centered (13), scaled (13)
```

```
## Resampling: Cross-Validated (10 fold, repeated 5 times)
```

```
## Summary of sample sizes: 630, 631, 630, 630, 630, 630, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	C	RMSE	Rsquared	MAE
##	0.25	0.2642019	0.7652614	0.1843101
##	0.50	0.2527741	0.7781828	0.1773106
##	1.00	0.2453811	0.7870477	0.1733668
##	2.00	0.2416795	0.7906762	0.1732208
##	4.00	0.2439411	0.7860016	0.1768360
##	8.00	0.2485535	0.7779967	0.1818068
##	16.00	0.2591541	0.7596403	0.1906157
##	32.00	0.2737784	0.7355134	0.2024909
##	64.00	0.2887308	0.7118227	0.2144476
##	128.00	0.3011335	0.6912139	0.2247847

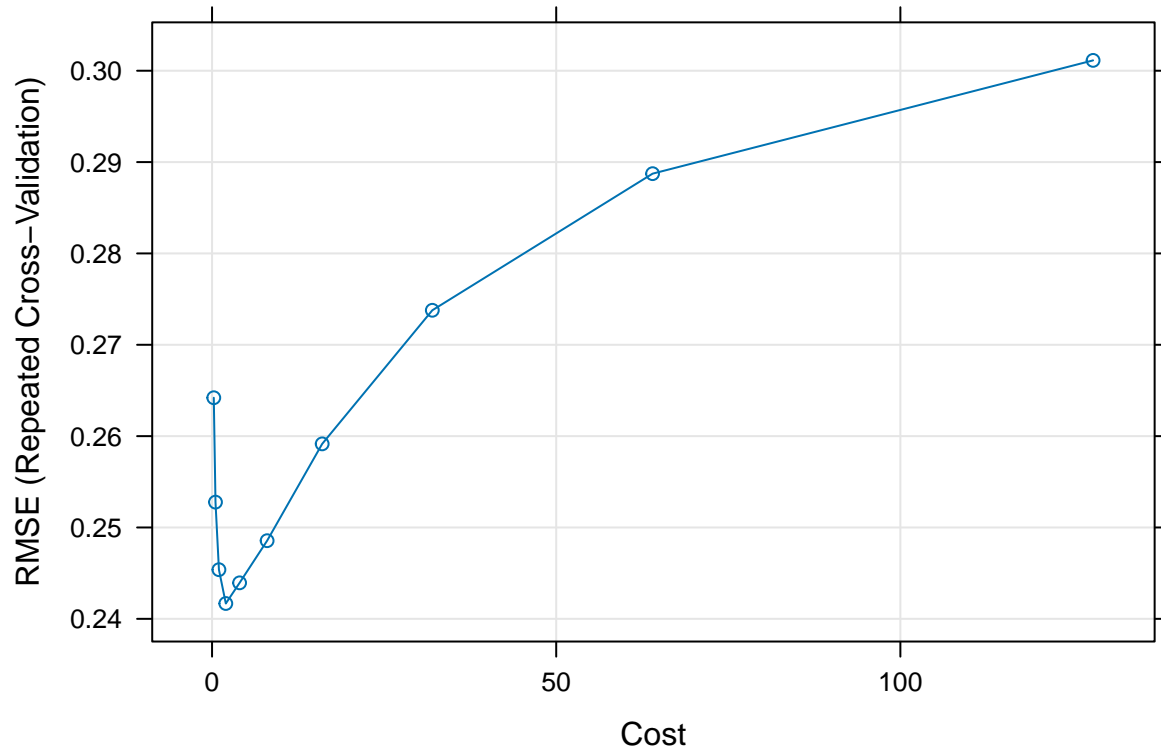
```
##
```

```
## Tuning parameter 'sigma' was held constant at a value of 0.05706539
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were sigma = 0.05706539 and C = 2.
```

```
plot(svmTuned)
```



```
#predicting on testing
svm_Pred <- predict(svmTuned, newdata = hp_test)
postResample(pred = svm_Pred, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.2573199 0.7715487 0.1850124
```

```
#####predicting on training
svm_Pred2 <- predict(svmTuned, newdata = hp_train)
postResample(pred = svm_Pred2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.1724535 0.8953050 0.1119898
```

```
###MARS###
marsGrid <- expand.grid(.degree = 1:2, .nprune = 2:10)

marsTuned <- train(hp_train, price_train,
  method = "earth",
  # Explicitly declare the candidate models to test
  tuneGrid = marsGrid,
  trControl = control,
  preProc = c("center", "scale"))
```

```
## Loading required package: earth
```

```
## Loading required package: Formula

## Loading required package: plotmo

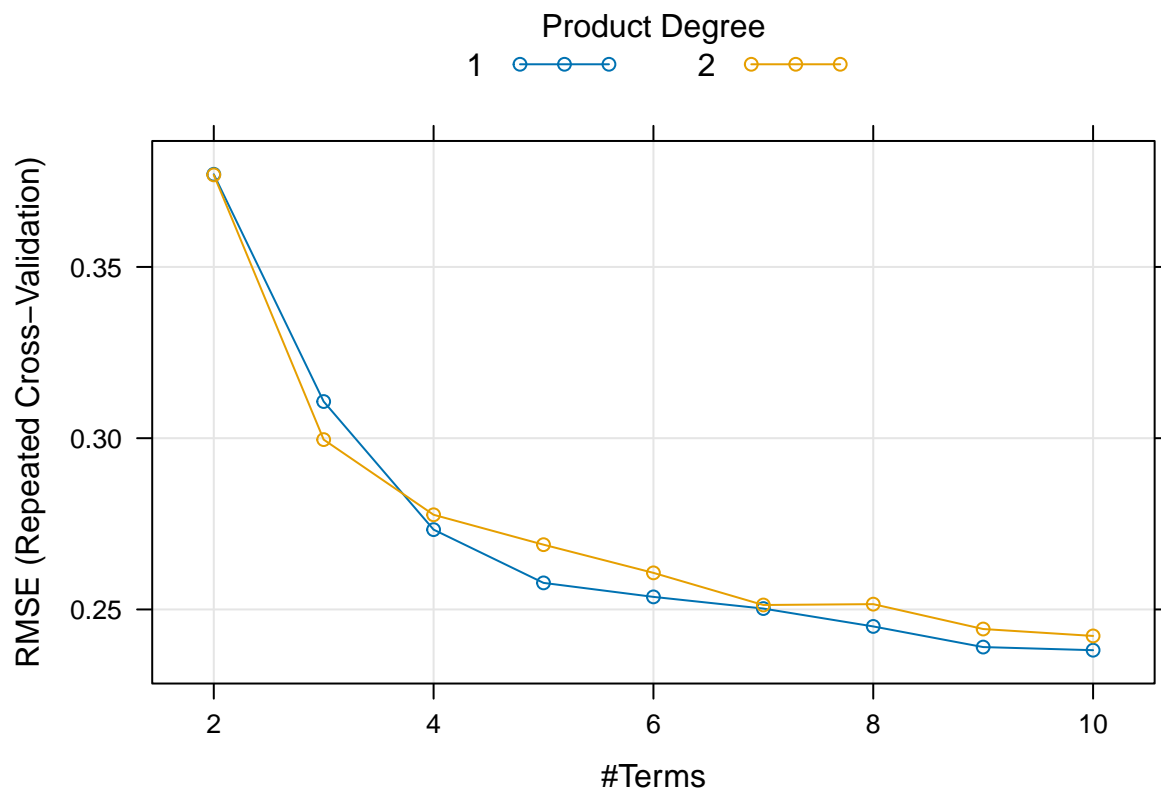
## Loading required package: plotrix

## Loading required package: TeachingDemos
```

```
marsTuned
```

```
## Multivariate Adaptive Regression Spline
##
## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 631, 630, 631, 630, 629, 630, ...
## Resampling results across tuning parameters:
##
## degree nprune RMSE Rsquared MAE
## 1 2 0.3770576 0.4869119 0.3014677
## 1 3 0.3107070 0.6541342 0.2323831
## 1 4 0.2733069 0.7325886 0.2053682
## 1 5 0.2577547 0.7626523 0.1924864
## 1 6 0.2536697 0.7701263 0.1892628
## 1 7 0.2502617 0.7765076 0.1859455
## 1 8 0.2450444 0.7856212 0.1817108
## 1 9 0.2390153 0.7955651 0.1774155
## 1 10 0.2381025 0.7971292 0.1768117
## 2 2 0.3768749 0.4880086 0.3014053
## 2 3 0.2995890 0.6772812 0.2219382
## 2 4 0.2776388 0.7234517 0.2086327
## 2 5 0.2689113 0.7405377 0.2024942
## 2 6 0.2606856 0.7566485 0.1953263
## 2 7 0.2513025 0.7733294 0.1875974
## 2 8 0.2515470 0.7739203 0.1865617
## 2 9 0.2442830 0.7869092 0.1819096
## 2 10 0.2422668 0.7914077 0.1796245
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 10 and degree = 1.
```

```
plot(marsTuned)
```



*#predicting on the testing set*

```
mars_Pred <- predict(marsTuned, hp_test)
postResample(pred = mars_Pred, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.2717943 0.7405466 0.1985179
```

*##predicting on the training set*

```
mars_Pred2 <- predict(marsTuned, hp_train)
postResample(pred = mars_Pred2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.2313384 0.8070711 0.1703204
```

#####

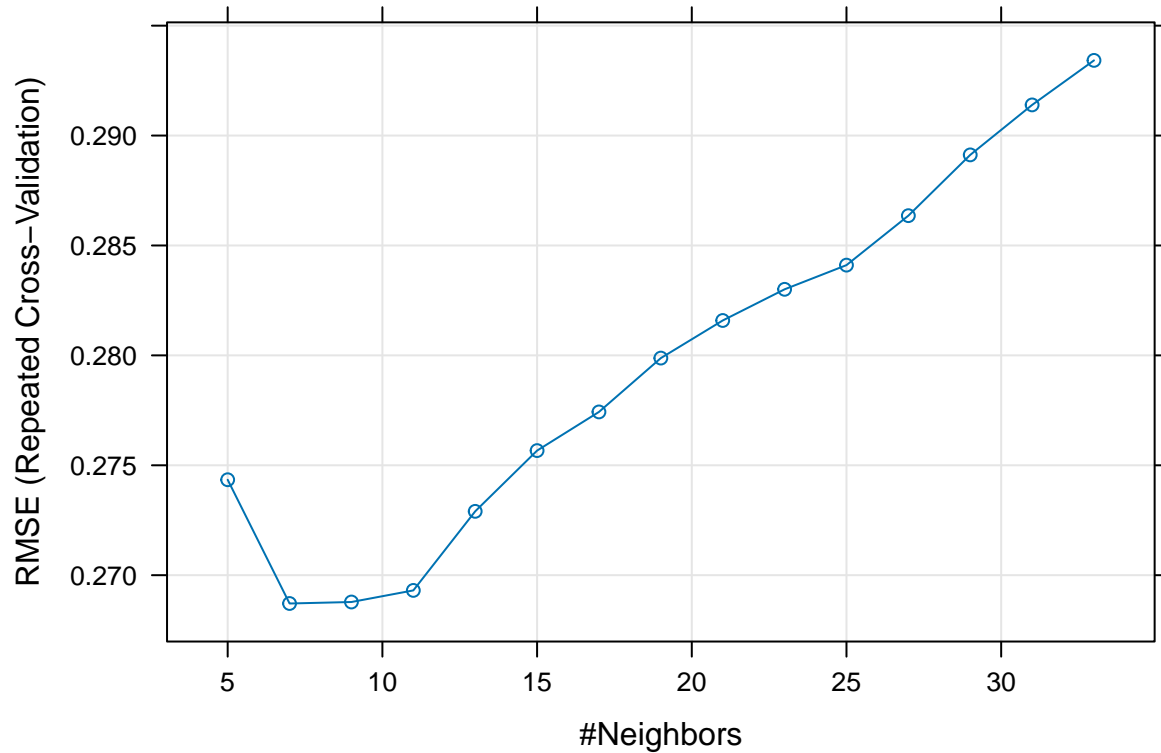
###KNN###

```
knnTuned <- train(hp_train,
                  price_train,
                  method = "knn",
                  preProc = c("center", "scale"),
                  trControl = control,
                  tuneLength = 15)
```

knnTuned

```
## k-Nearest Neighbors
##
## 700 samples
## 13 predictor
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 631, 630, 630, 629, 631, 629, ...
## Resampling results across tuning parameters:
##
##   k    RMSE      Rsquared    MAE
##   5  0.2743424  0.7326776  0.1993332
##   7  0.2687121  0.7484982  0.1933011
##   9  0.2687799  0.7514694  0.1928847
##  11  0.2693075  0.7546081  0.1937398
##  13  0.2729046  0.7514835  0.1960295
##  15  0.2756683  0.7486809  0.1973130
##  17  0.2774268  0.7488425  0.1989028
##  19  0.2798771  0.7477018  0.2011146
##  21  0.2815872  0.7482426  0.2021891
##  23  0.2830055  0.7493502  0.2027897
##  25  0.2841063  0.7505927  0.2032161
##  27  0.2863556  0.7497497  0.2049116
##  29  0.2891212  0.7469890  0.2066425
##  31  0.2913948  0.7448141  0.2083133
##  33  0.2934210  0.7436622  0.2099352
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 7.
```

```
plot(knnTuned)
```



```
#predictin on the test set
knn_Pred <- predict(knnTuned, hp_test)
postResample(pred = knn_Pred, price_test)
```

```
##      RMSE  Rsquared      MAE
## 0.2930136 0.7168625 0.2145057
```

```
#####predicting on the training
knn_Pred2 <- predict(knnTuned, hp_train)
postResample(pred = knn_Pred2, price_train)
```

```
##      RMSE  Rsquared      MAE
## 0.2310558 0.8166257 0.1660358
```